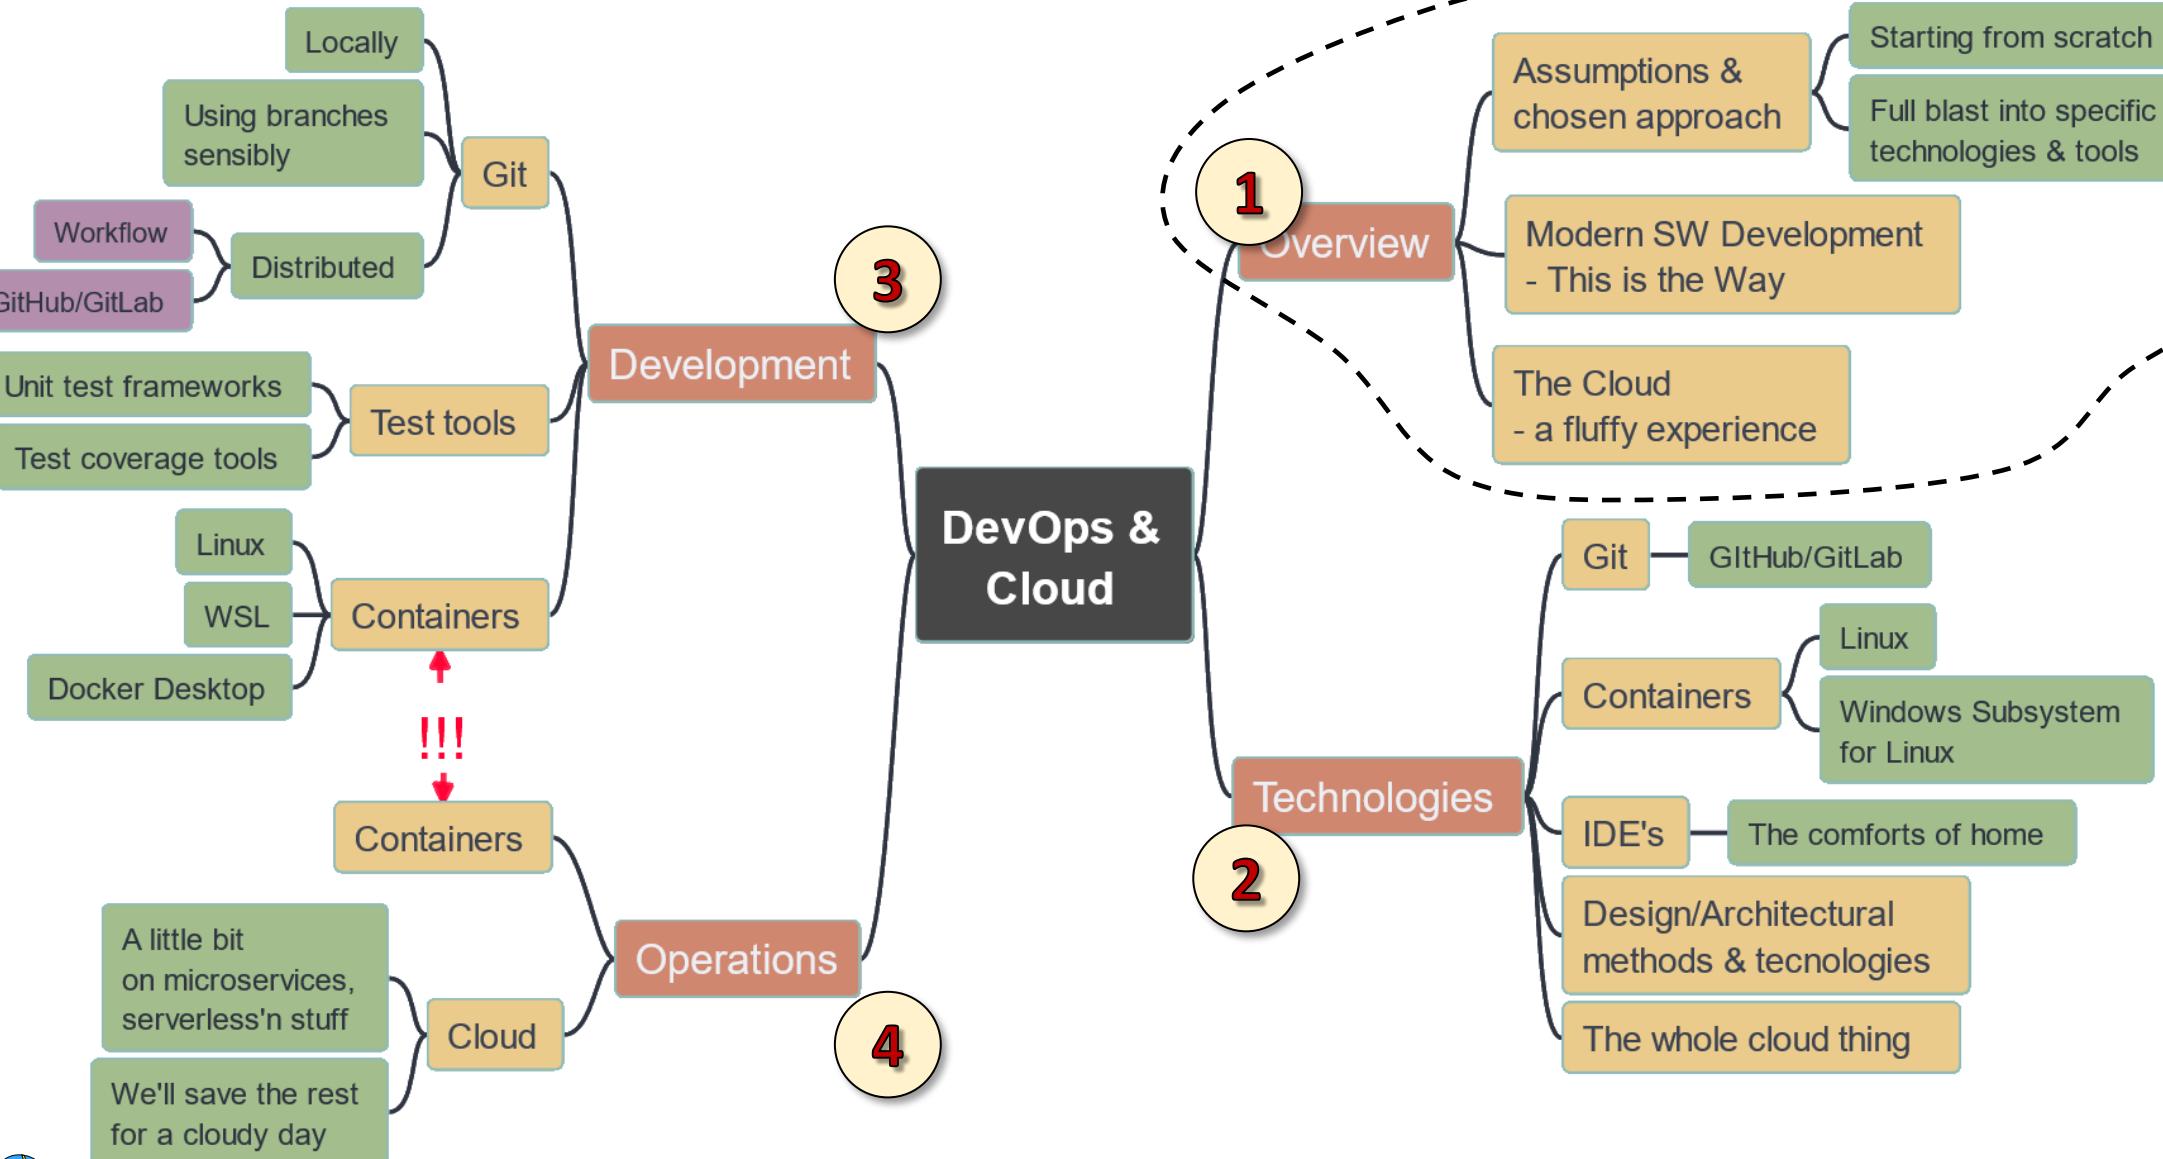


# DevOps & Cloud

- an introduction

# Overview – mind the map



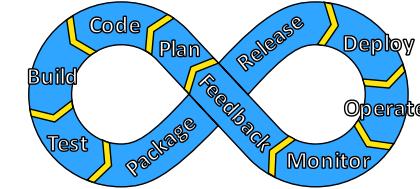


# About this mini-course

- Assumptions on my part
  - None – so I may regurgitate some stuff you already know (but this does no harm and it's important anyway, so there ....)
- Approach
  - Some very general introductions
  - Then headfirst into the tools of the trade
  - To get you up to speed thinking and talking about, and using, these tools and technologies as fast as possible
  - To allow you to navigate the terms and concepts (which are not in any way difficult but somewhat fuzzy at times)



# DevOps – what is it? (1)



- The short answer: That's the way modern sw development is done
- The slightly longer answer:
  - Modern sw development (SWD) recognizes that SWD is much more than simply writing (good code)
  - We need to look at the full SWD lifecycle, from conception, over development and operations, to the day the sw is scrapped – good sw development processes must facilitate all phases
  - So, the “word” DevOps = **D**evelopment And **O**perations indicates methods and technologies that cover the full lifecycle:
    - Development = Specify, design, implement, test, package and release software
    - Operations = Putting the sw into operation, maintain, configure, update and monitor software
  - The “DevOps loop” indicates that the parts (Dev & Ops) feed off and into each other





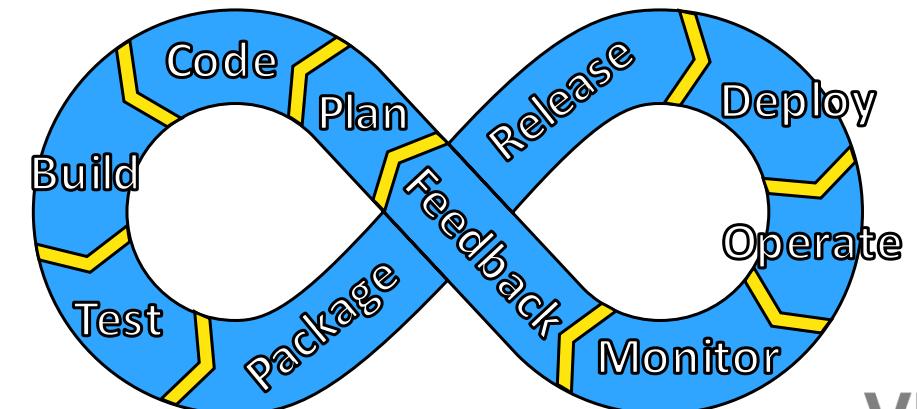
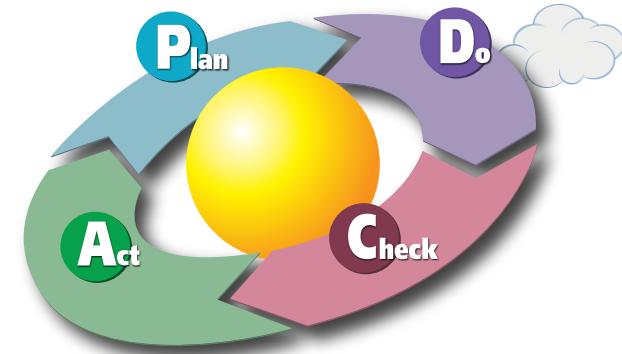
# DevOps – what is it? (2)

- The whole Ops part obviously depends on the context and type of sw – here's two different scenarios:
  1. Sw that is fully developed and operated internally in a company. The sw is put into immediate use each time it is updated.
  2. Sw that it is developed as an off-the-shelf product. The sw gets updated regularly and official releases made once in while.
- Scenario 1 has the closest coupling between dev & ops, and may benefit the most from DevOps methods and technologies. The ops part provides direct feedback to developers about the quality and usability of the se
- But even the developers in scenario 2 should think about (some of) the operations, especially the update part (nobody likes sw that nukes everything because it gets updated) but also the feedback from users!



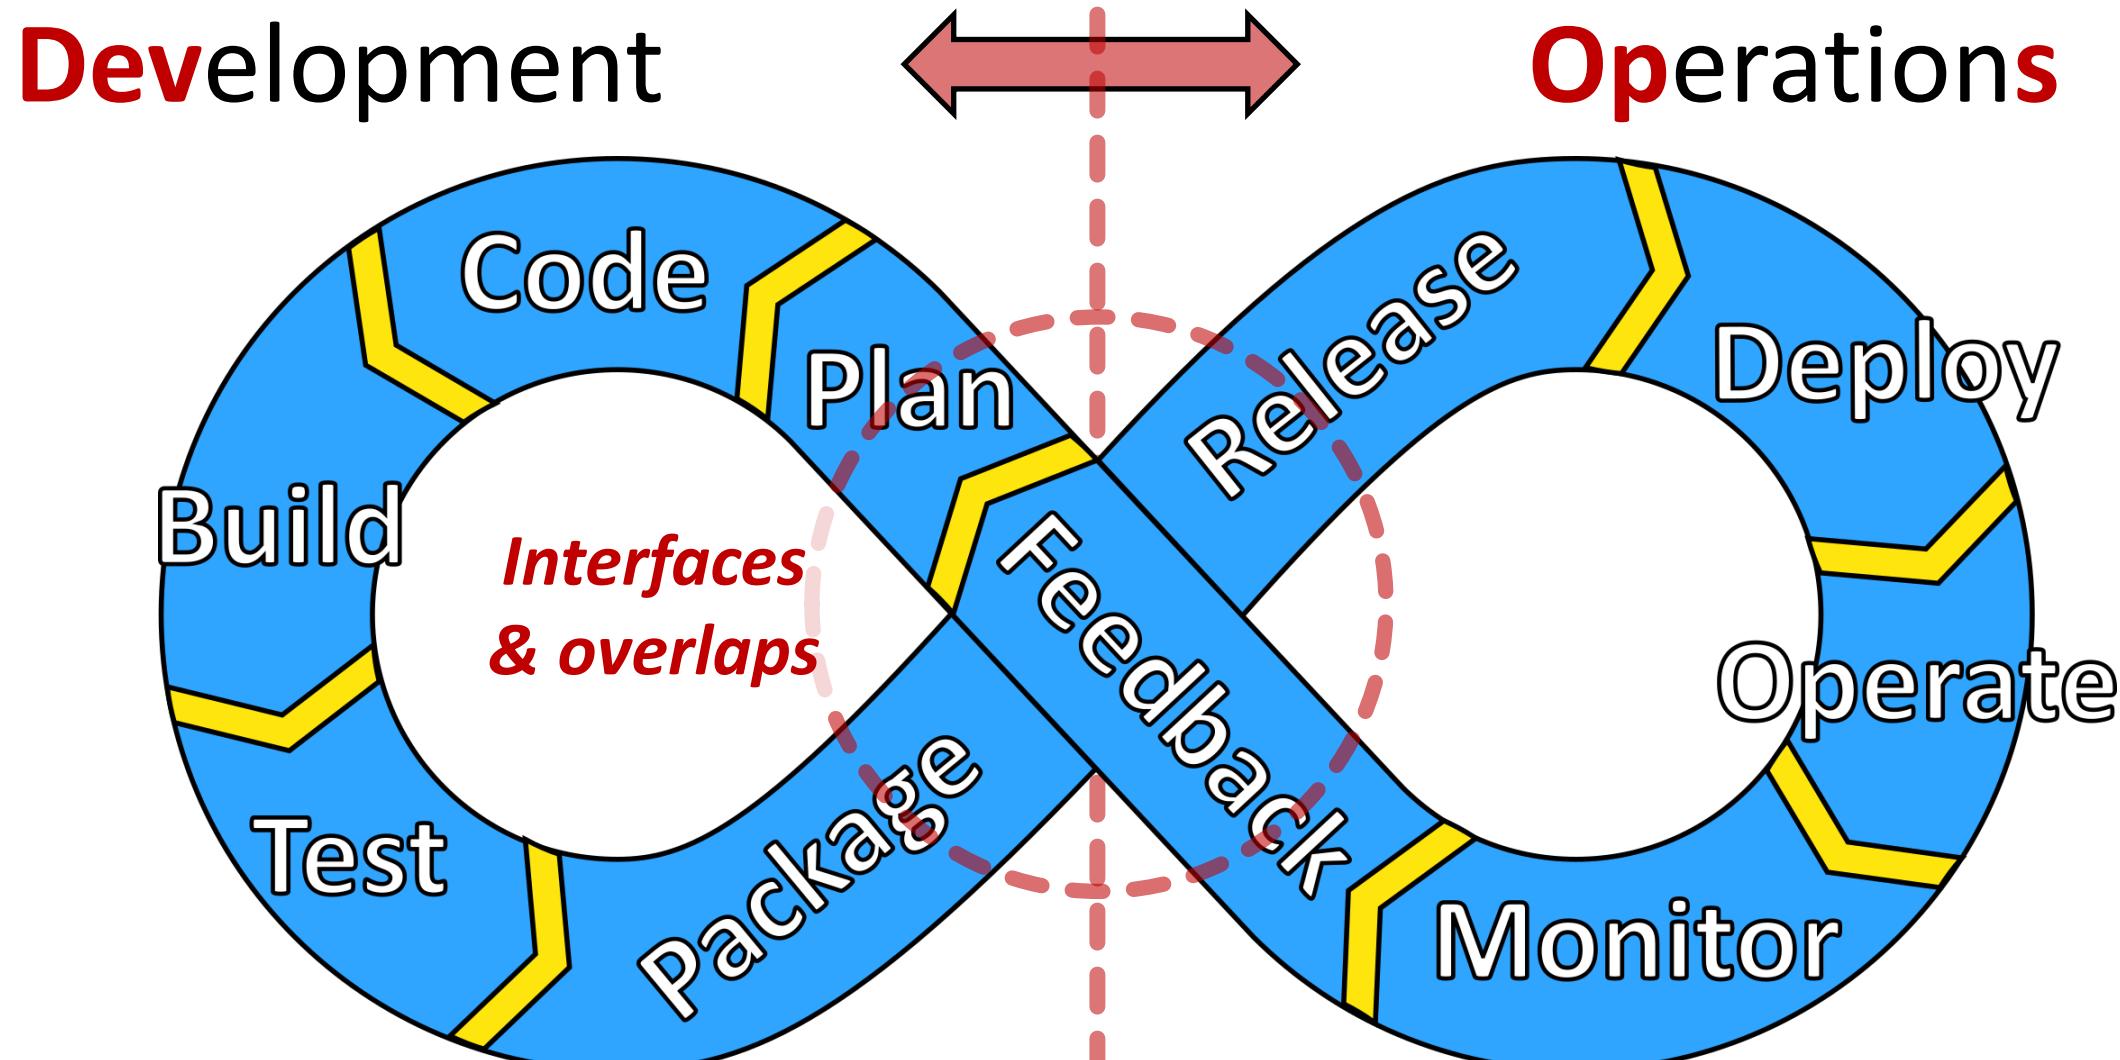
# DevOps – what is it? (3)

- Modern development processes promotes agile and iterative sw development where the sw is built incrementally in many small specify-design-implement/test-release cycles (aka p-d-c-a)
- We want to continuously integrate new features in the code (**CI** = continuous integration) and make the result immediately available for use (**CD** = continuous delivery/deployment)
- So, we need to be efficient at doing all this specify/design/ implement/ test/release/package stuff
- To do that we have to think about the whole process and not only about writing some code:





# The DevOps loop



# Cloud – why's that relevant? (1)



- Cloud is just a common designation for a set of technologies
- Cloud is mostly associated with the Ops-part of DevOps (deploying, running and monitoring applications), but many newer development tools also make good use of cloud technologies
- The “cloud” thing indicates that these technologies are commonly realized on computing resources that are owned by another party, and that a company pay the cloud owner for using those resources based on how much you use them (on-demand, pay-per-use)
- But there's absolutely nothing that prevents a company from using “cloud” technologies on computers bought and owned by the company itself

# Cloud – why's that relevant? (2)



- Using a “real cloud” relieves you from buying, maintaining and operating computer hardware which can be costly (money, man-hours, power, cooling etc.)
- But after the initial cloud hype there’s been a trend the last few years of companies moving back to self-owned hardware because it may actually be cheaper, even though you then have to pay for hw, salaries and power
- A mix of on-premise resources and real cloud (aka hybrid cloud) is also becoming common
  - You can use the cloud to quickly prototype new setups and then move it internally when things work
  - You can use your own stack and use the cloud to help handle peak loads





# Cloud – why's that relevant? (3)

- Stuff like GDPR also muddies the water when it comes to using the “real” cloud – it may prevent certain customers from using a company’s product if the customers have no control of the physical whereabouts of data
- So, all in all, the notion of “cloud” incorporates several ways of using these technologies
  1. Full-on real cloud from providers like Amazon, Azure and Google
  2. Hybrid cloud combining self-owned resources with external “real cloud”
  3. Fully internal “cloud” running 100% on self-owned resources
- But that does not diminish the value of the technologies involved, such as virtualization and network-based storage - the cloud thing is really about making efficient use of computer resources by sharing them between many applications in a flexible and dynamic way





# What is DevOps, here's 3 takes - a bit fuzzy, right?

1. Used as *a set of practices and tools*, **DevOps** integrates and automates the work of software development (Dev) and IT operations (Ops) as *a means for improving and shortening the systems development life cycle*.
2. Other than it being a cross-functional combination (and a portmanteau) of the terms and concepts for "development" and "operations", academics and practitioners have not developed a universal definition for the term "**DevOps**". Most often, **DevOps** is characterized by *key principles: shared ownership, workflow automation, and rapid feedback*.
3. "*a set of practices intended to reduce the time between committing a change to a system and the change being placed into normal production, while ensuring high quality*"





# DevOps in summary – our take

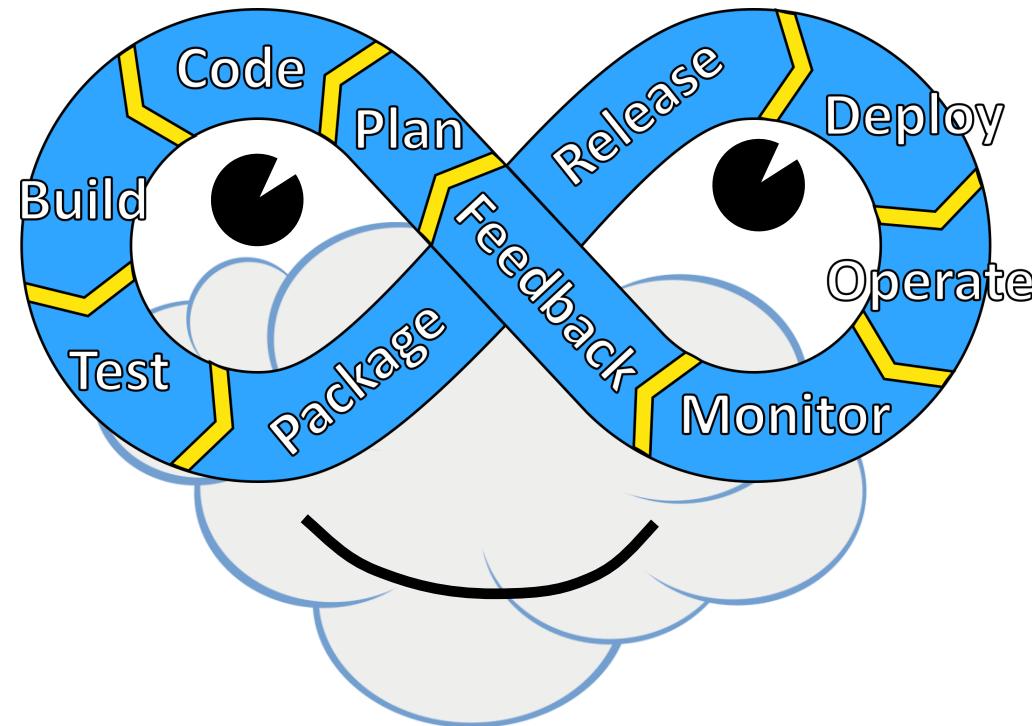
- 1) A *way of working*, creating *smart workflows*, that *reduces the time* that goes from deciding and planning a feature (or bugfix) until
  - That feature is implemented, fully tested and *integrated in our code base*
  - A new version of our system is *in normal operation*
- 2) Any *tool* or *method* that helps us achieving the above by
  - *automating* development and operational workflows, across *the full software development life-cycle*
  - *supporting the simultaneous work* of multiple developers working on several lines of development
  - *bridging the technical differences* between the development environment and the operational environment
  - *facilitates deployment* of new releases of software





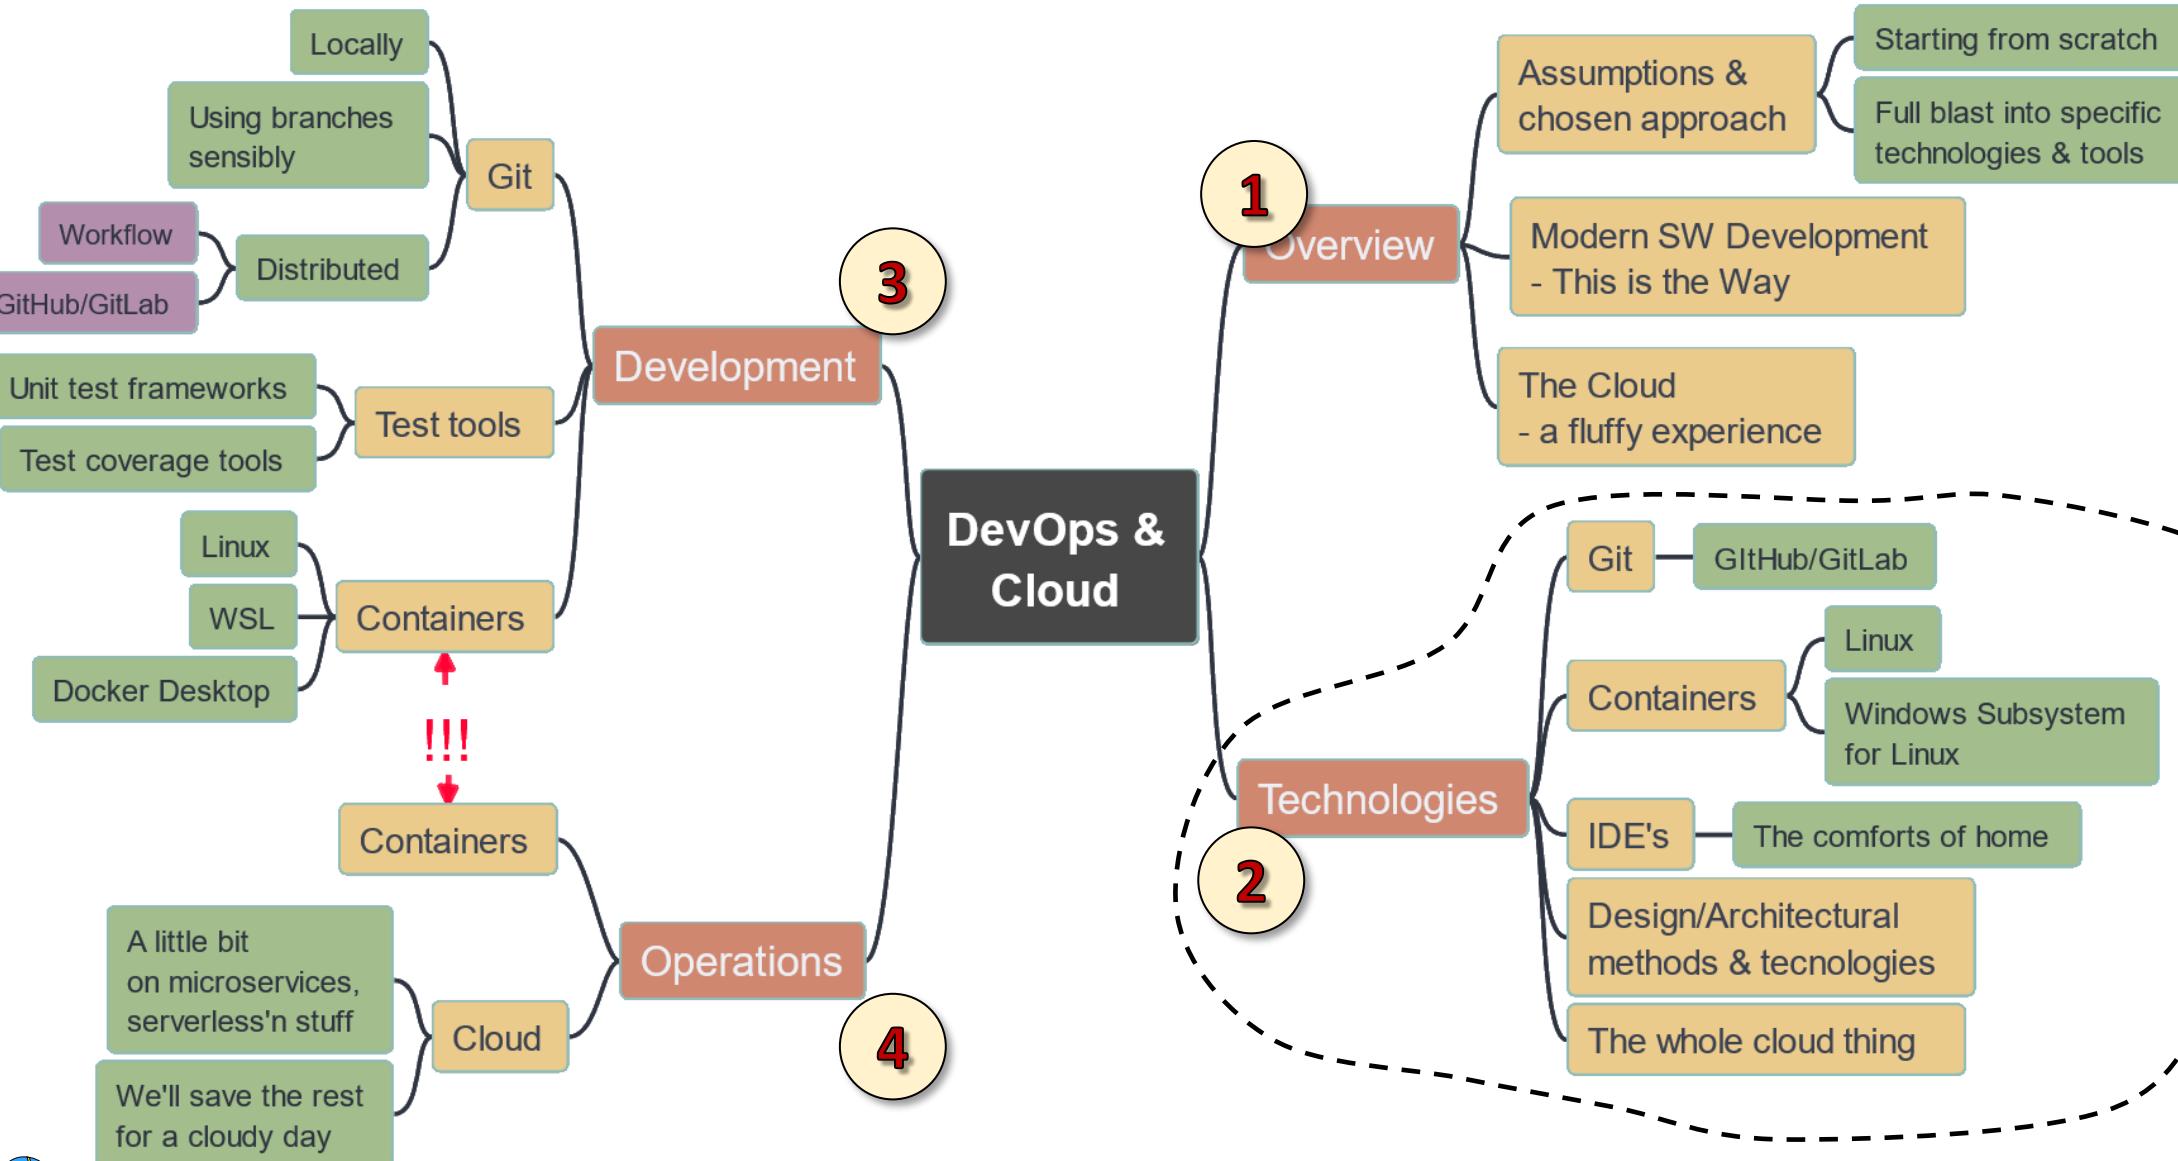
# DevOps and Cloud - BFF

Using the definition from the previous slide, we can see why DevOps and Cloud concepts are BFF – Cloud technology can really help us achieve the DevOps goals of minimizing the systems development life cycle time





# Overview – mind this map



# Technologies – quick appetizer .... (1)

- ***Git*** (VCS - version control system) – because that's the stuff that ties things together and forms the core of our workflow
- ***Linux*** (often in the form of WSL (Windows Subsystem for Linux)) – because Linux is ubiquitous and (for DevOps/Cloud) has one special feature that is reeeeeaaaally helpful, namely
- ***Containers*** (a form of virtualization) – because this tech makes it possible for us to fully control the runtime environment of our applications in a cheap, neatly wrapped packet





# Technologies – quick appetizer .... (2)

- ***Design/architectural*** methods – because they make it possible to make sw that is both scalable, robust and easy to move around (even while an application is running). Examples are:
  - Microservices
  - Serverless
- ***IDE's*** (such as IntelliJ and Visual Studio Code) – because they are equipped with features that works with containers which make it possible to do all sorts of neat tricks
- ***Github/Gitlab*** – because they facilitate both distributed development and automation of builds, testing, packaging and releases





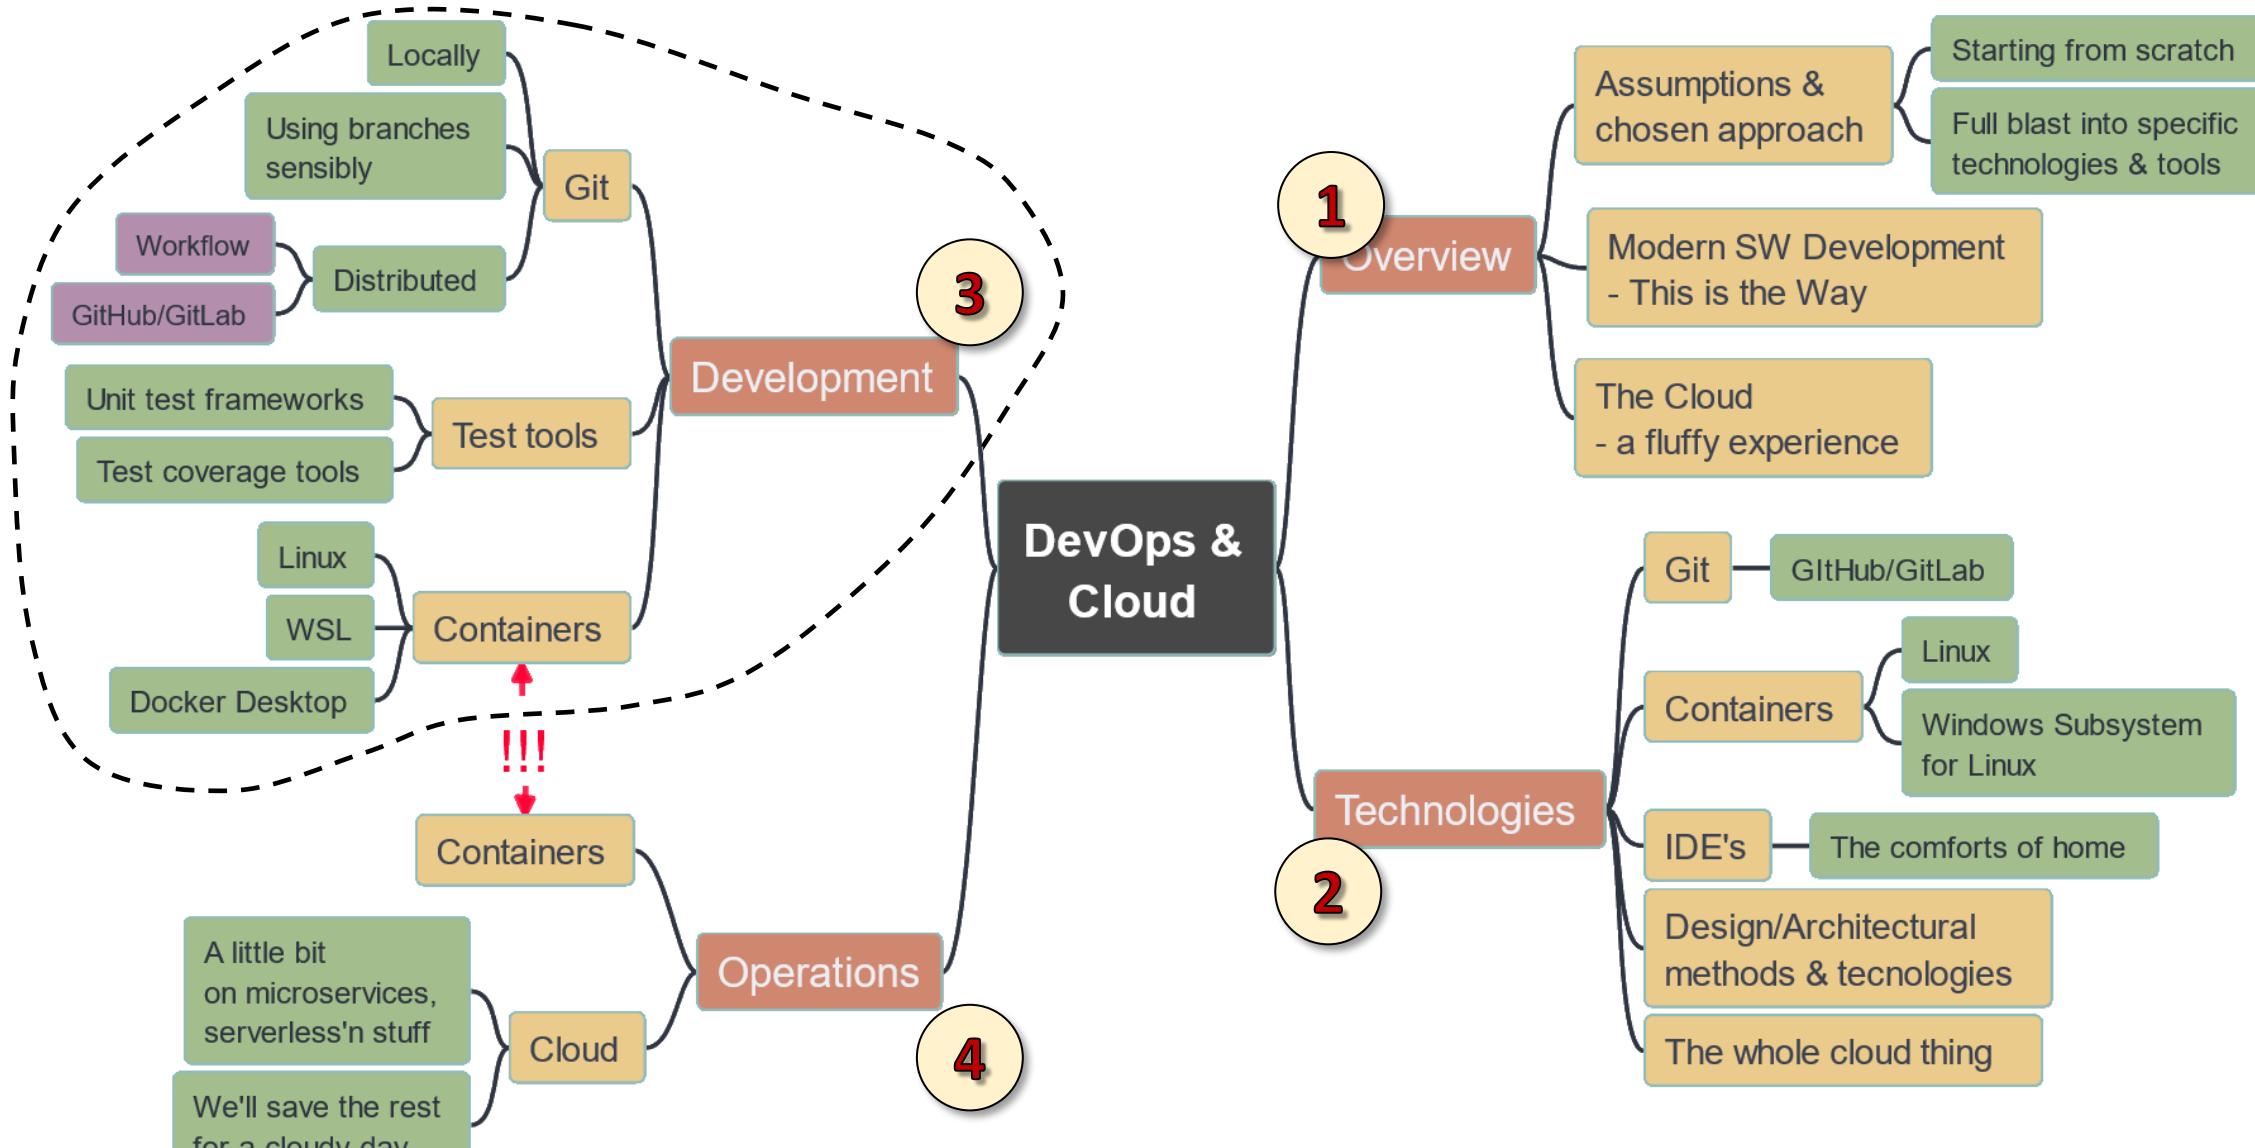
# Technologies – quick appetizer .... (3)

- This whole Cloud thing – Amazon, Google, Azure, OpenStack, etc.
- The Cloud providers more or less provide the same services and resources
  - CPU cycles (virtual machines running, Linux or Windows, commissioned in seconds) and computer RAM
  - Persistent storage (in several categories, getting more expensive the faster you want)
  - Network (bandwidth, DNS mappings, and what have you not)
  - Lots of “premade” appliances (databases (SQL or NoSql), webservers etc.) for us to use or customize





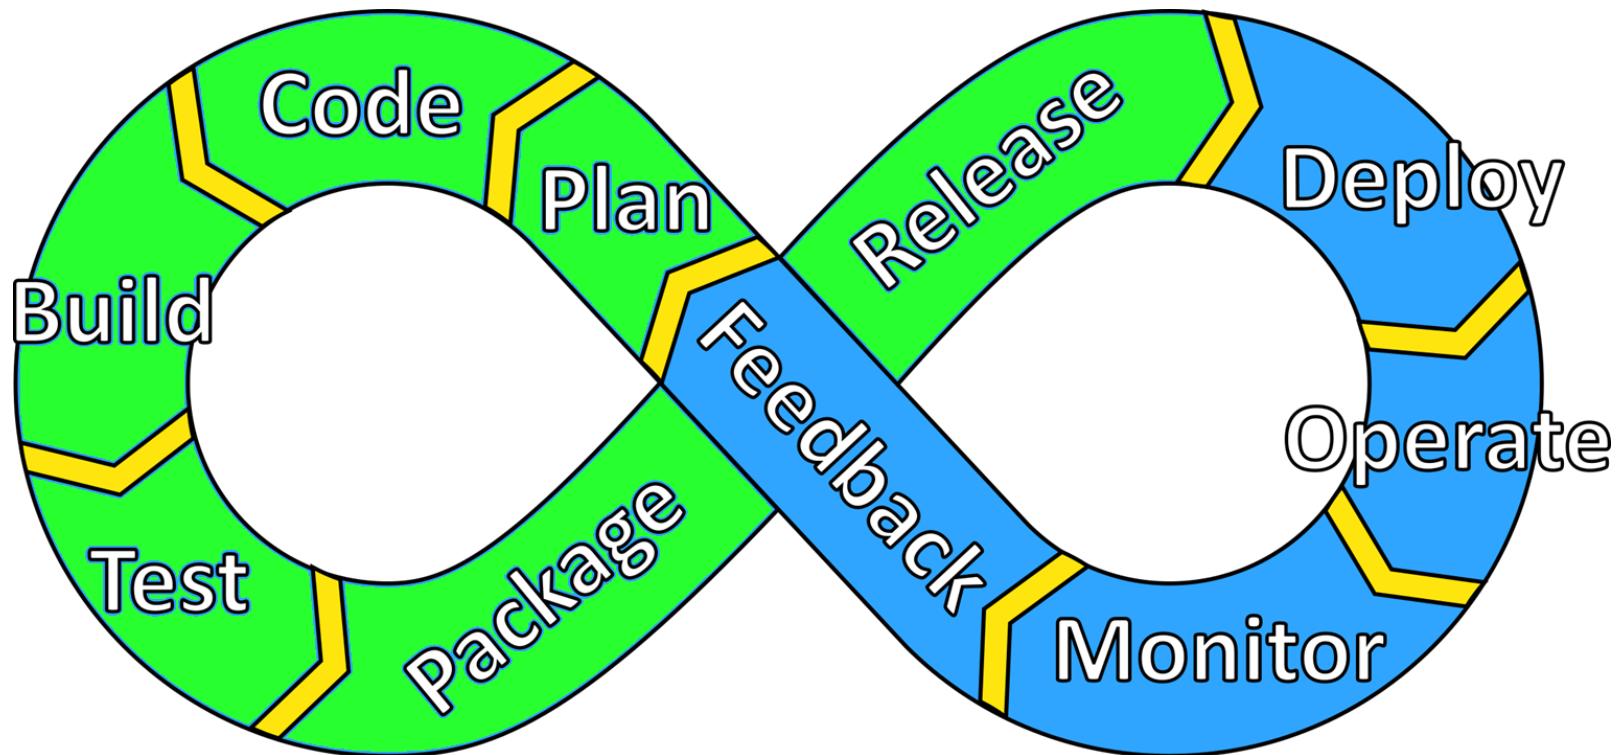
# Overview – mind this map



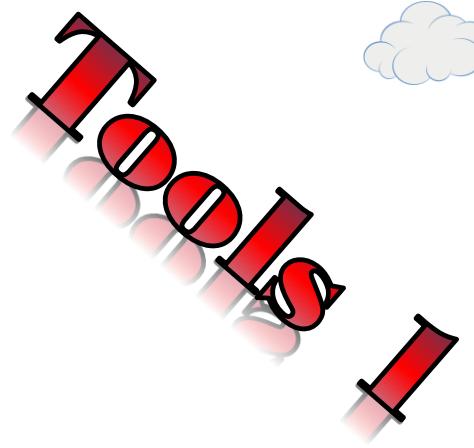


# The Dev-part of DevOps – tools & technologies

- So let's start out with a look at (some of) the tools/technologies we can use for DevOps'ing the Dev part ....



# Git basics – using Git sensibly



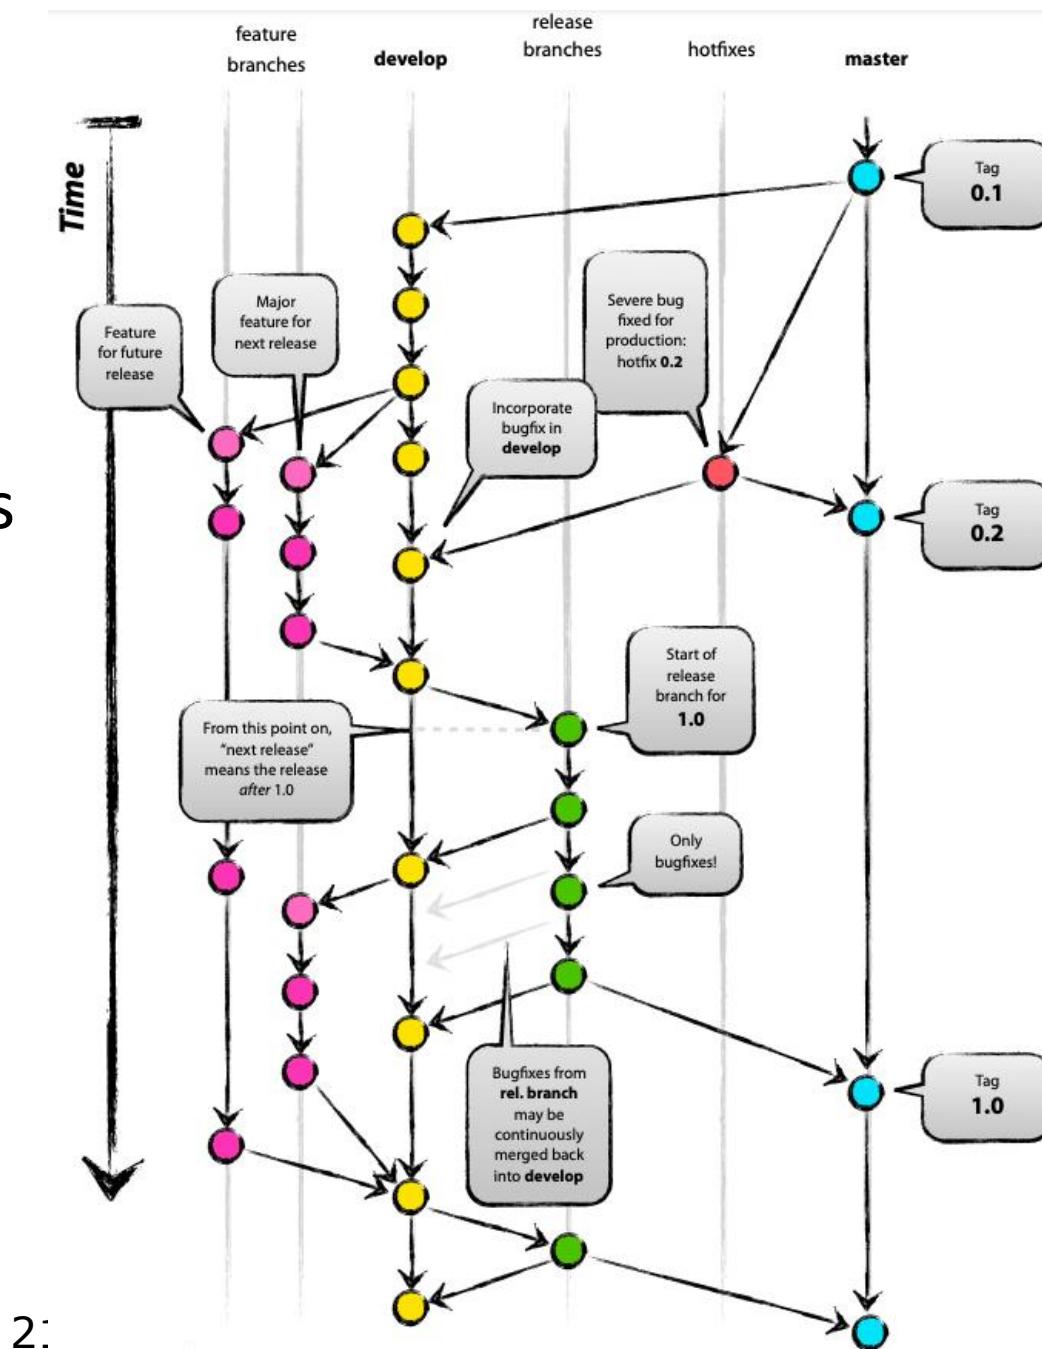
- We have to think about Git in two (but related) contexts:
  - 1) **Locally** (and that's all we discussed so far)
  - 2) **Distributed** (we'll get into more details on that later)
- Local use of Git
  - Something you do from a command prompt with git commands
  - Or something you do from your IDE of choice
  - We need to be able to work locally with branches, tags etc. before we start sharing code
- Distributed use of Git
  - This is where GitHub/GitLab etc. comes into play
  - Here we must work smart with branches, pull requests, branch protection rules etc.





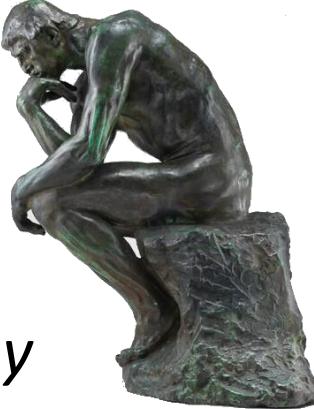
# GitFlow – overkill!

- Googling around, many stumbles across GitFlow
- A model for using Git branches in a large project with many participants and parallel ongoing lines of work (features, bug fixes etc.)
- ***Don't use that for your project!***
- GitFlow has its raison d'etre and a role to play in some circumstances, but ***it's not*** for everybody



# Thinking branches “globally” and locally ...

main... fixes...  
features...



- First a few observations:
  1. We can do whatever we want locally, in our own copy of a central repo.  
*So: as long as we don't try to push our local branches to the repo, they are invisible to everybody else using (cloned copies of) the central repo.*
  2. Branches in a central repo (i.e. GitHub, GitLab etc.) are visible to everybody that clones that repo.  
*So: all users of a shared repo needs to agree to a set of rules about how the “public” branches are used.*
  3. One fixed and crucial rule for branches in a central repo is *that all changes (commits/merges) to the main development branch must be done via pull requests* in the central repo.





# To *push* local changes, we need to do a *pull* request

To push a change on a local branch to a global (= shared) branch we must perform three steps:

- **First push the local branch, warts and all, to the central repo (Github).**  
That just creates/updates a copy of the local branch on Github – it has no effect on any other branches
- **Then create a pull request from the shared branch where you want the updates to end up.**  
The shared branch should be setup disallow direct pushes and require approval of pull requests – so this step does not change the shared either
- **Finally, when the pull request has been reviewed and approved, merge the pull request into the shared branch**

First at this point do the changes hit the shared branch and become visible to others



# Thinking branches “globally” and locally ...

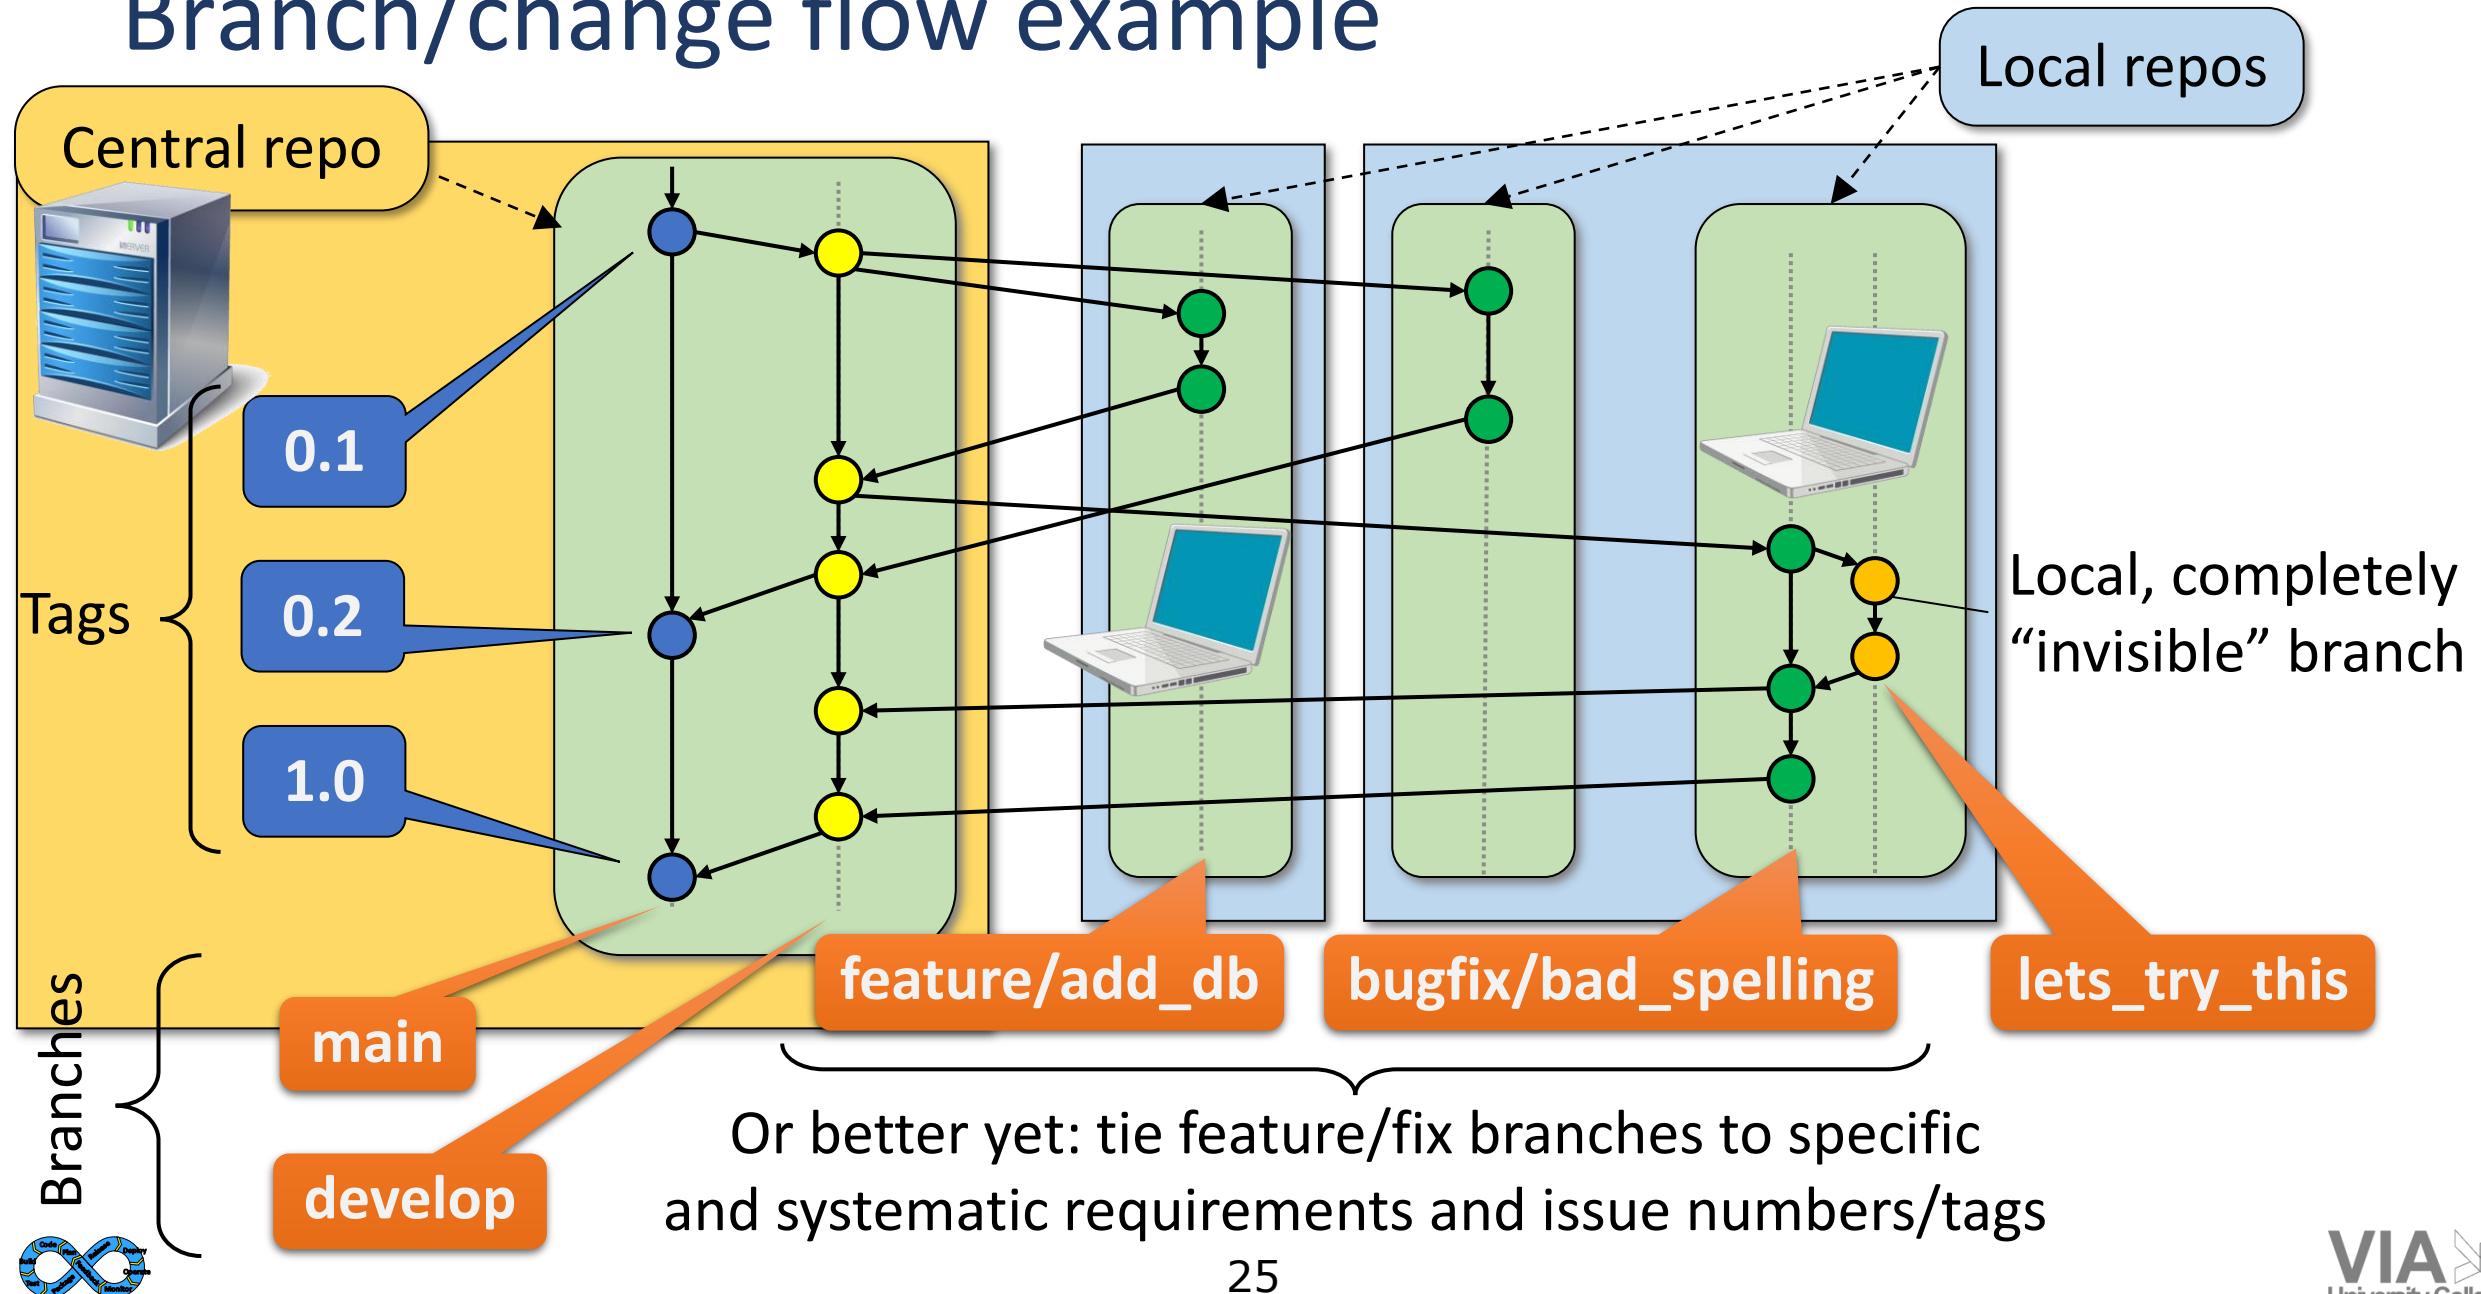


- As long as we keep these things in mind, we can do pretty much whatever we want
- Here's a few ***SUGGESTIONS*** .... (kind of a minimum viable rule set)
  - Globally (in GitHub/Gitlab)
    - Use a single branch to track your “releases” (could be **main**)
    - Use a designated branch for integration of features and fixes (let’s call it **develop**)
    - Snapshot **develop** to **main** everytime you want to release (and add a suitable tag)
  - Locally (in your cloned copy of the central repo)
    - Create and use a new branch, branching off from **develop**, for every new feature or bug fix you want to make
    - Use well-defined names for feature branches (feature/<description>) and bug fix branches (bugfix/<description>)
    - Once the local feature/fix is done, push the branch to GitHub/GitLab and create a pull req.





# Branch/change flow example





# Git – good git manners (how best to use Git)

To make the most of Git (for yourself and your team members), use these rules:

- 1) Use branches all the time, use them sensibly and name them well (example: <https://dev.to/varbsan/a-simplified-convention-for-naming-branches-and-commits-in-git-il4>)
- 2) Make many frequent and small commits, rather than a seldom and few large commits – use selective staging to break down larger changes into smaller commits
- 3) Make each commit be about one thing, and 1 thing only and use sensible rules for the commit messages (example: <https://cbea.ms/git-commit/>)
- 4) Do NOT check in derived/generated files (such as compiled stuff) – EVER!
- 5) Do NOT check in config files that change screen colors, keyboard shortcuts or any other personal preferences – your teammates will (rightly) kill you!
- 6) Setup, use and maintain one or more **.gitignore** files from the very start



# Good house keeping ...

- Delete branches when they are no longer needed. You can always recreate the branch from the last commit if needed
- This is a must in the central repo – otherwise the old branches with completed features and fixed bugs will seep into local repos over time and pollute the whole thing
- You can do what you want locally, but your local repo can easily drown in a mountain of lint in the shape of old, unused and stale branches – so cleanup



- Good test tools and test habits are critical for keeping the code base in good shape
- Using good test tools (especially for unit testing), and using them right and all the time, is very much a part of a DevOps flow
- Automating tests not only saves time writing and maintaining the tests but saves TONS of time downstream
- Adding a test coverage tool adds a fantastic objective view into the state of the code base
  - Which parts have been tested, and which parts not?
  - How does test coverage evolve over time? Is it steady/increasing/decreasing?

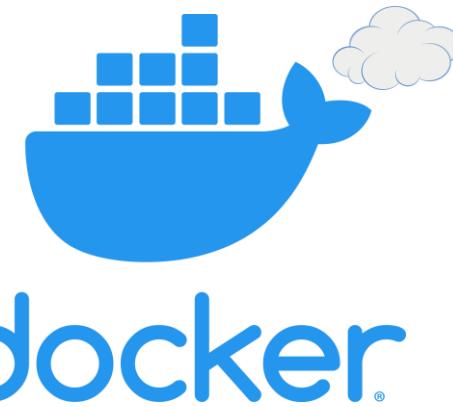


# How to enforce workflow rules on GitHub

- We have 2 essential tools in GitHub that we can use to enforce workflow rules that help us keep our code base in good shape:
  - Always buildable
  - All tests running
  - Always runnable
- These 2 tools are:
  - Branch protection rules
  - GitHub actions – which really should be called GitHub workflows
- A small example <https://github.com/DaFessor/devops-demo>
- Branch protection, workflows, test and coverage

= code that's in good shape



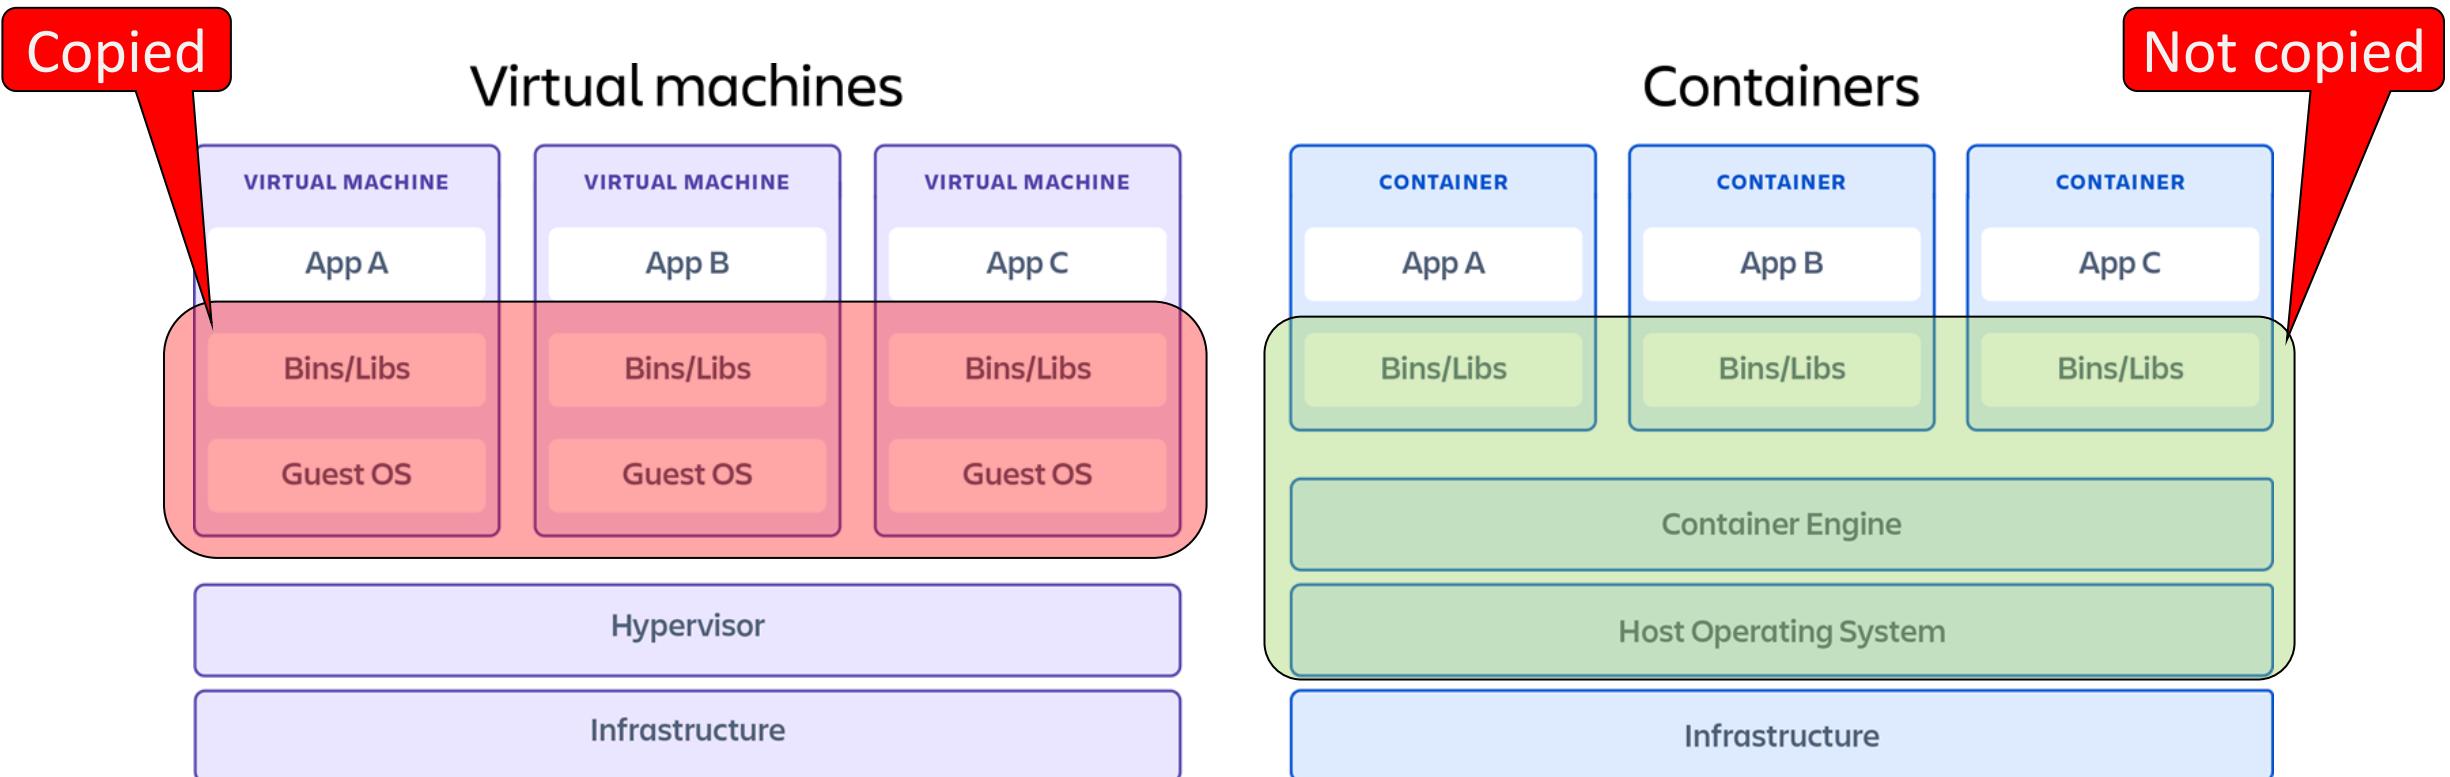


- Docker is the company that popularized Linux containers
- A container is a kind of “virtual machine”, in the sense that apps running inside a container sees their own “private” machine
  - They have their own file system, nobody from outside the container can mess with files
  - No other processes are running
  - They see their own network (devices and network addresses)
  - All needed libraries have the exact right, well-defined version
  - An app may run with root/superuser rights
- Seen from the outside the container is a fully closed and sandboxed jail
  - Apps running inside the container cannot see or modify files, processes or access the network outside the box



# Containers vs traditional VM's

- Containers start & stop really quick since they don't have to simulate a full machine w. hw and everything, VM's have to do a full boot/shutdown cycle
- Real VM's duplicate all stuff used across different VM's, containers don't





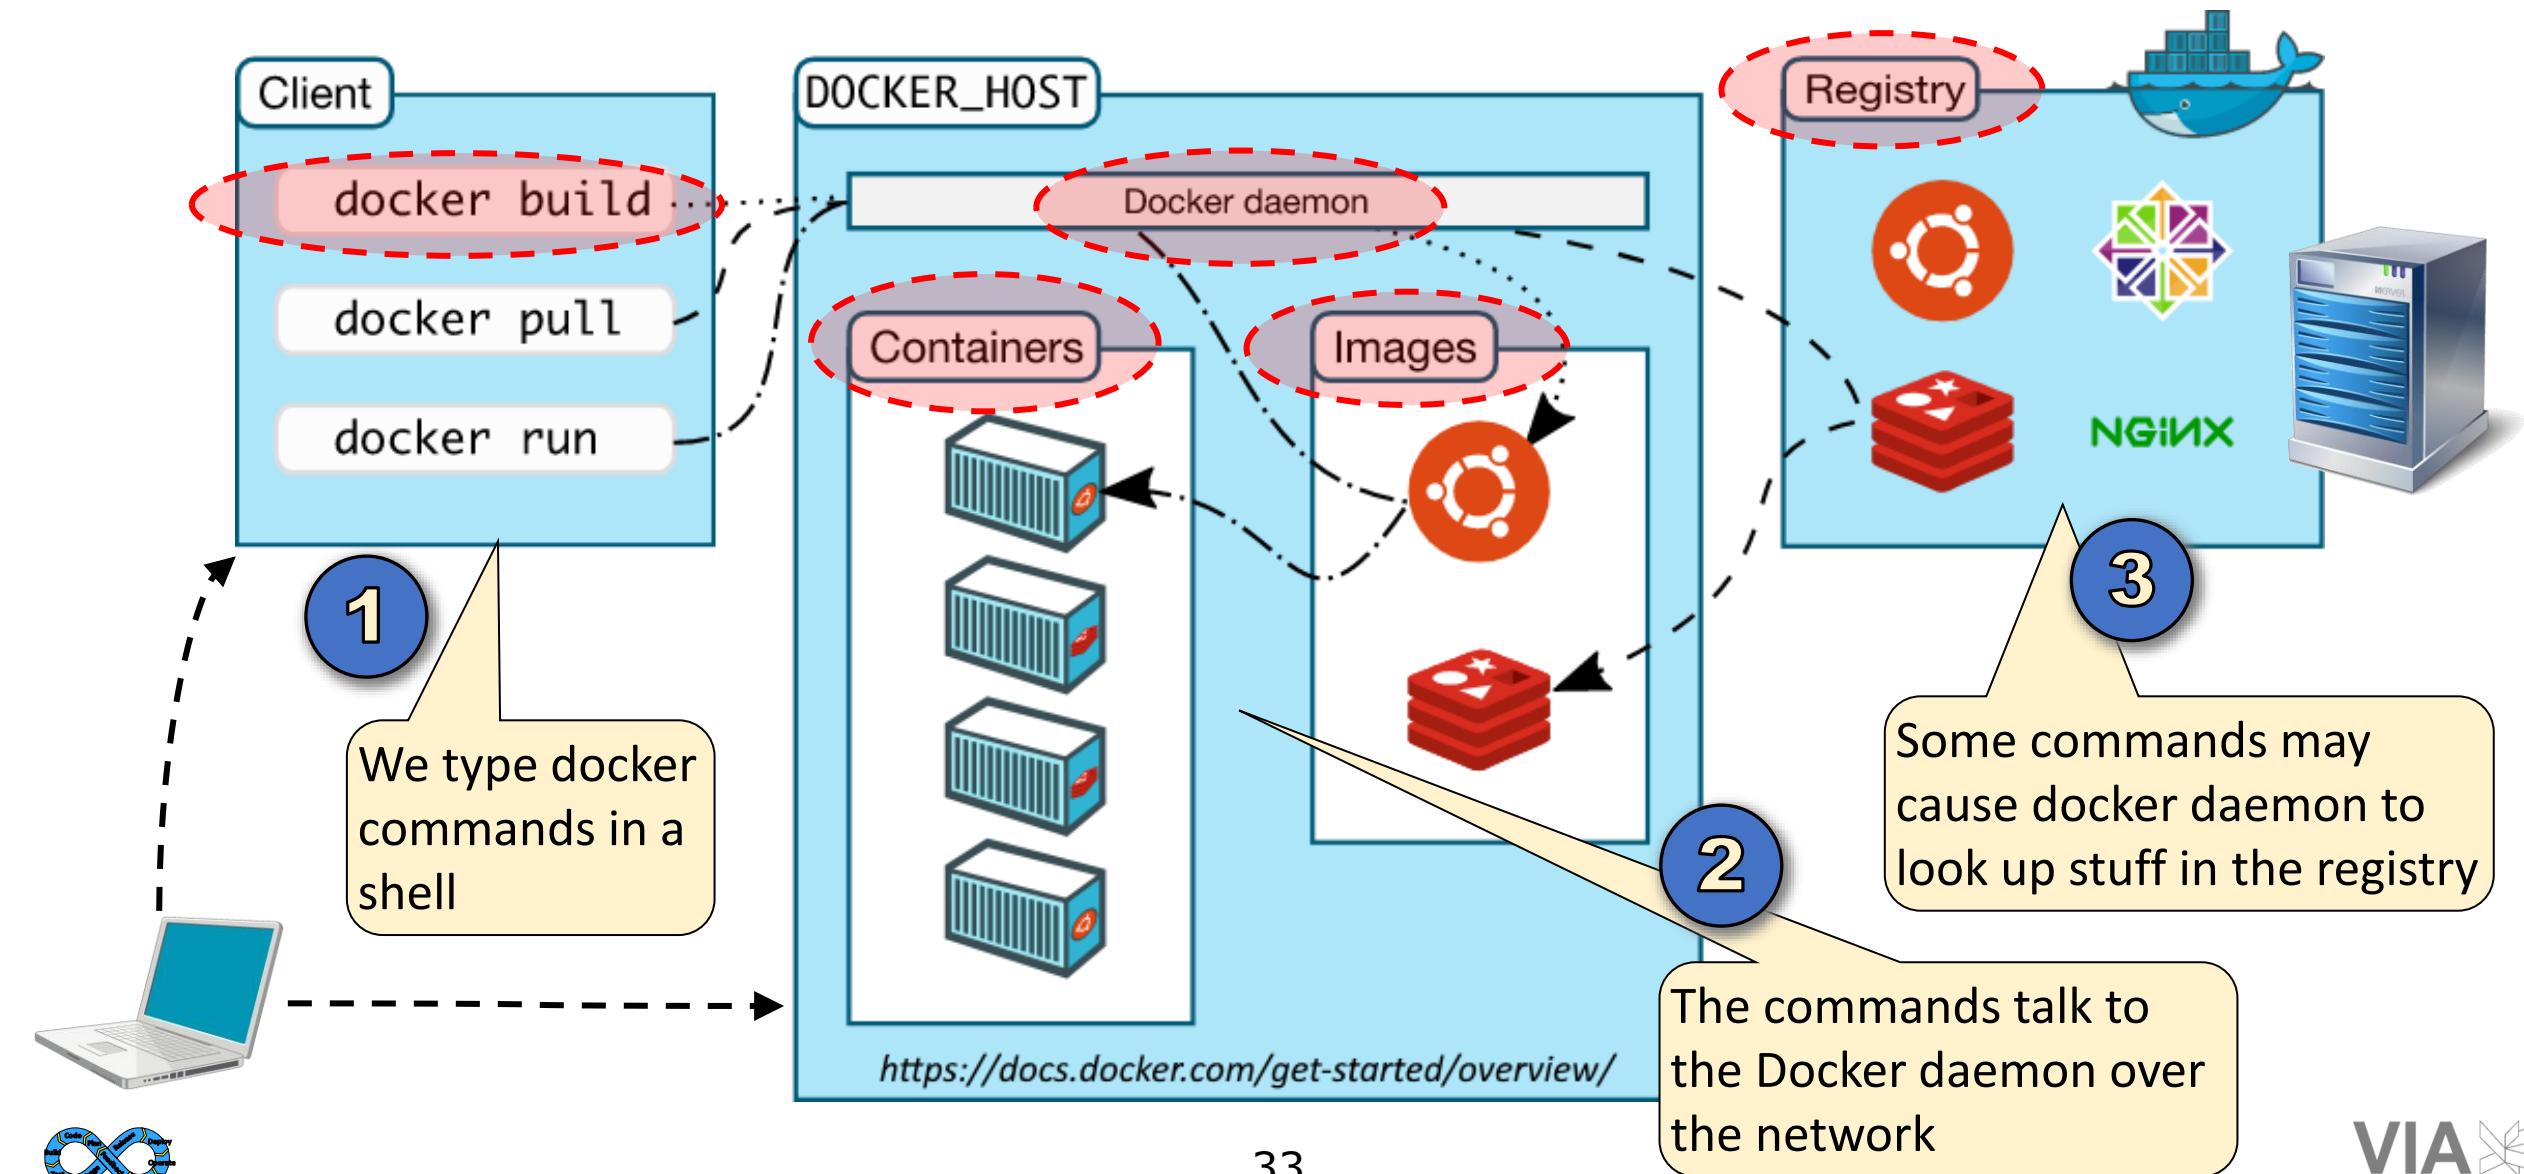
# Docker schmocker, Podman and containers

- Docker (the company) popularized containers, and we often use the two, Docker and containers, as if they are one and the same
- They are not – there are other providers of container technology on Linux, such as Redhat Linux that makes Podman
- So, informally Docker and Podman are two examples of *container management “engines”* allowing us to build, run and manage containers
- Container tech is standardized to an actually useful level, via the OCI (Open Container Initiative)
- Not all containers are the same type , but what we learn by working w. Docker works very well to other container types (e.g Podman)
- But henceforth we'll only talk about Docker containers





# How Docker works





# Docker images, containers and such ...

- **Client:** An application that wants to do Docker stuff. Example: A shell running in a terminal, where we manually execute Docker commands.
- **Dockerd:** The actual container management engine that runs as a service which the clients can talk to over a (usually local) network connection (i.e. client and docker host is often the same machine)
- **Registry:** One or more registries (“library”) of premade docker images ready to be downloaded and used
- **Image:** A read-only template with instructions for creating a Docker container, often based on another image with some additional customization.
- **Container:** A runnable instance of an image.
- **Volume:** A persistent “file system” that can be attached to a container.





# To work w. (Docker) containers we need Linux

- .... or a very recent version of macOS
- Docker Desktop also runs on Windows (handy) but it needs Windows Subsystem for Linux (WSL) to do handle the actual containerization
- WSL v2 runs a real Microsoft customized Linux kernel inside Windows
- Docker needs WSL2 to work

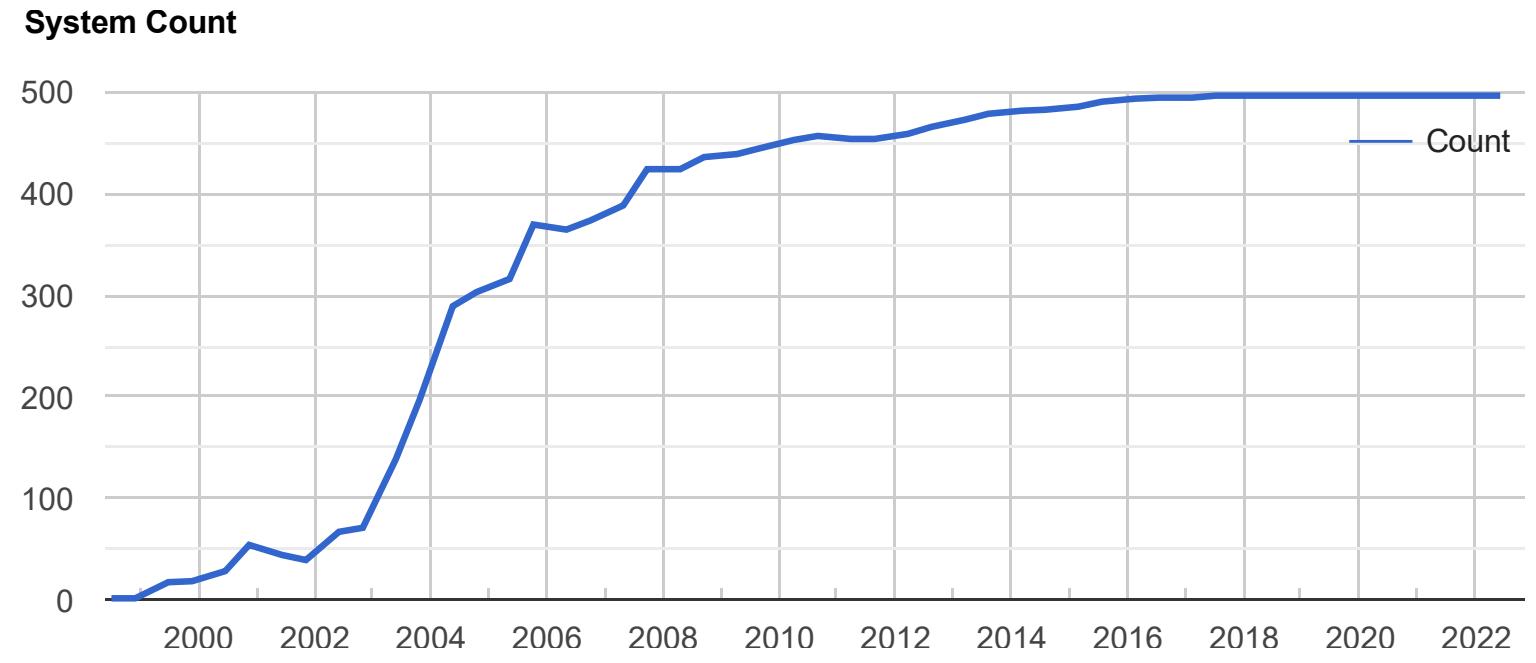




# Linux – from humble origins to today

- In **1991** Linus Torvalds, as a student of computer science at Helsinki University, wrote in an Internet news group:  
*Hello everybody out there using minix – I'm doing a (free) operating system (just a hobby, won't be big and professional like gnu) for 386(486) AT clones. This has been brewing since april, and is starting to get ready*

- A graph of how many of the most powerful 500 computers in the world that used Linux in november **2022**





# Linux usage

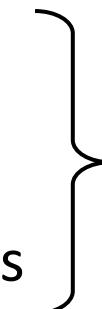
- In oct 2021 appr. 71% of all smart phones in the world used Android – which has Linux running under the surface
- On top of that comes all the routers, switches and other network devices, uses in automotive applications, smart home applications and so on ....
- In short, Linux is a **VERY WIDELY USED** open source operating system, that runs on all sorts of devices from usb-stick sized embedded devices to the largest computers in the world with literally thousands of cores.  
The frontrunner as of today('ish) has 8,730,112 cores, and delivers 7.94 EFlop/s (that's  $10^{18} = 7.94 \times 1 \text{ million} \times 1 \text{ million} \times 1 \text{ million}$  64-bit floating points ops per second – not shabby for a hobby project, eh? You can encode a LOT of cat videos VERY fast with that kind of power!





# Why has containers become so popular so fast?

- They are very use to manage using tools like Docker, Docker Compose and Kubernetes
- They are very fast to start up ( $\approx$ a few milliseconds on my humble laptop)
- They make it easy to run/deploy software together with all necessary dependencies as a single “runnable” entity
- They can be used everywhere we might want to run, test or build or our software:
  - On our development machines
  - In the pipelines of GitHub/GitLab etc.
  - On the operational target machine(s)/systems



***We can use the same container in all 3 situations!  
That has far-ranging implications for making things easier .....***



# Demo time

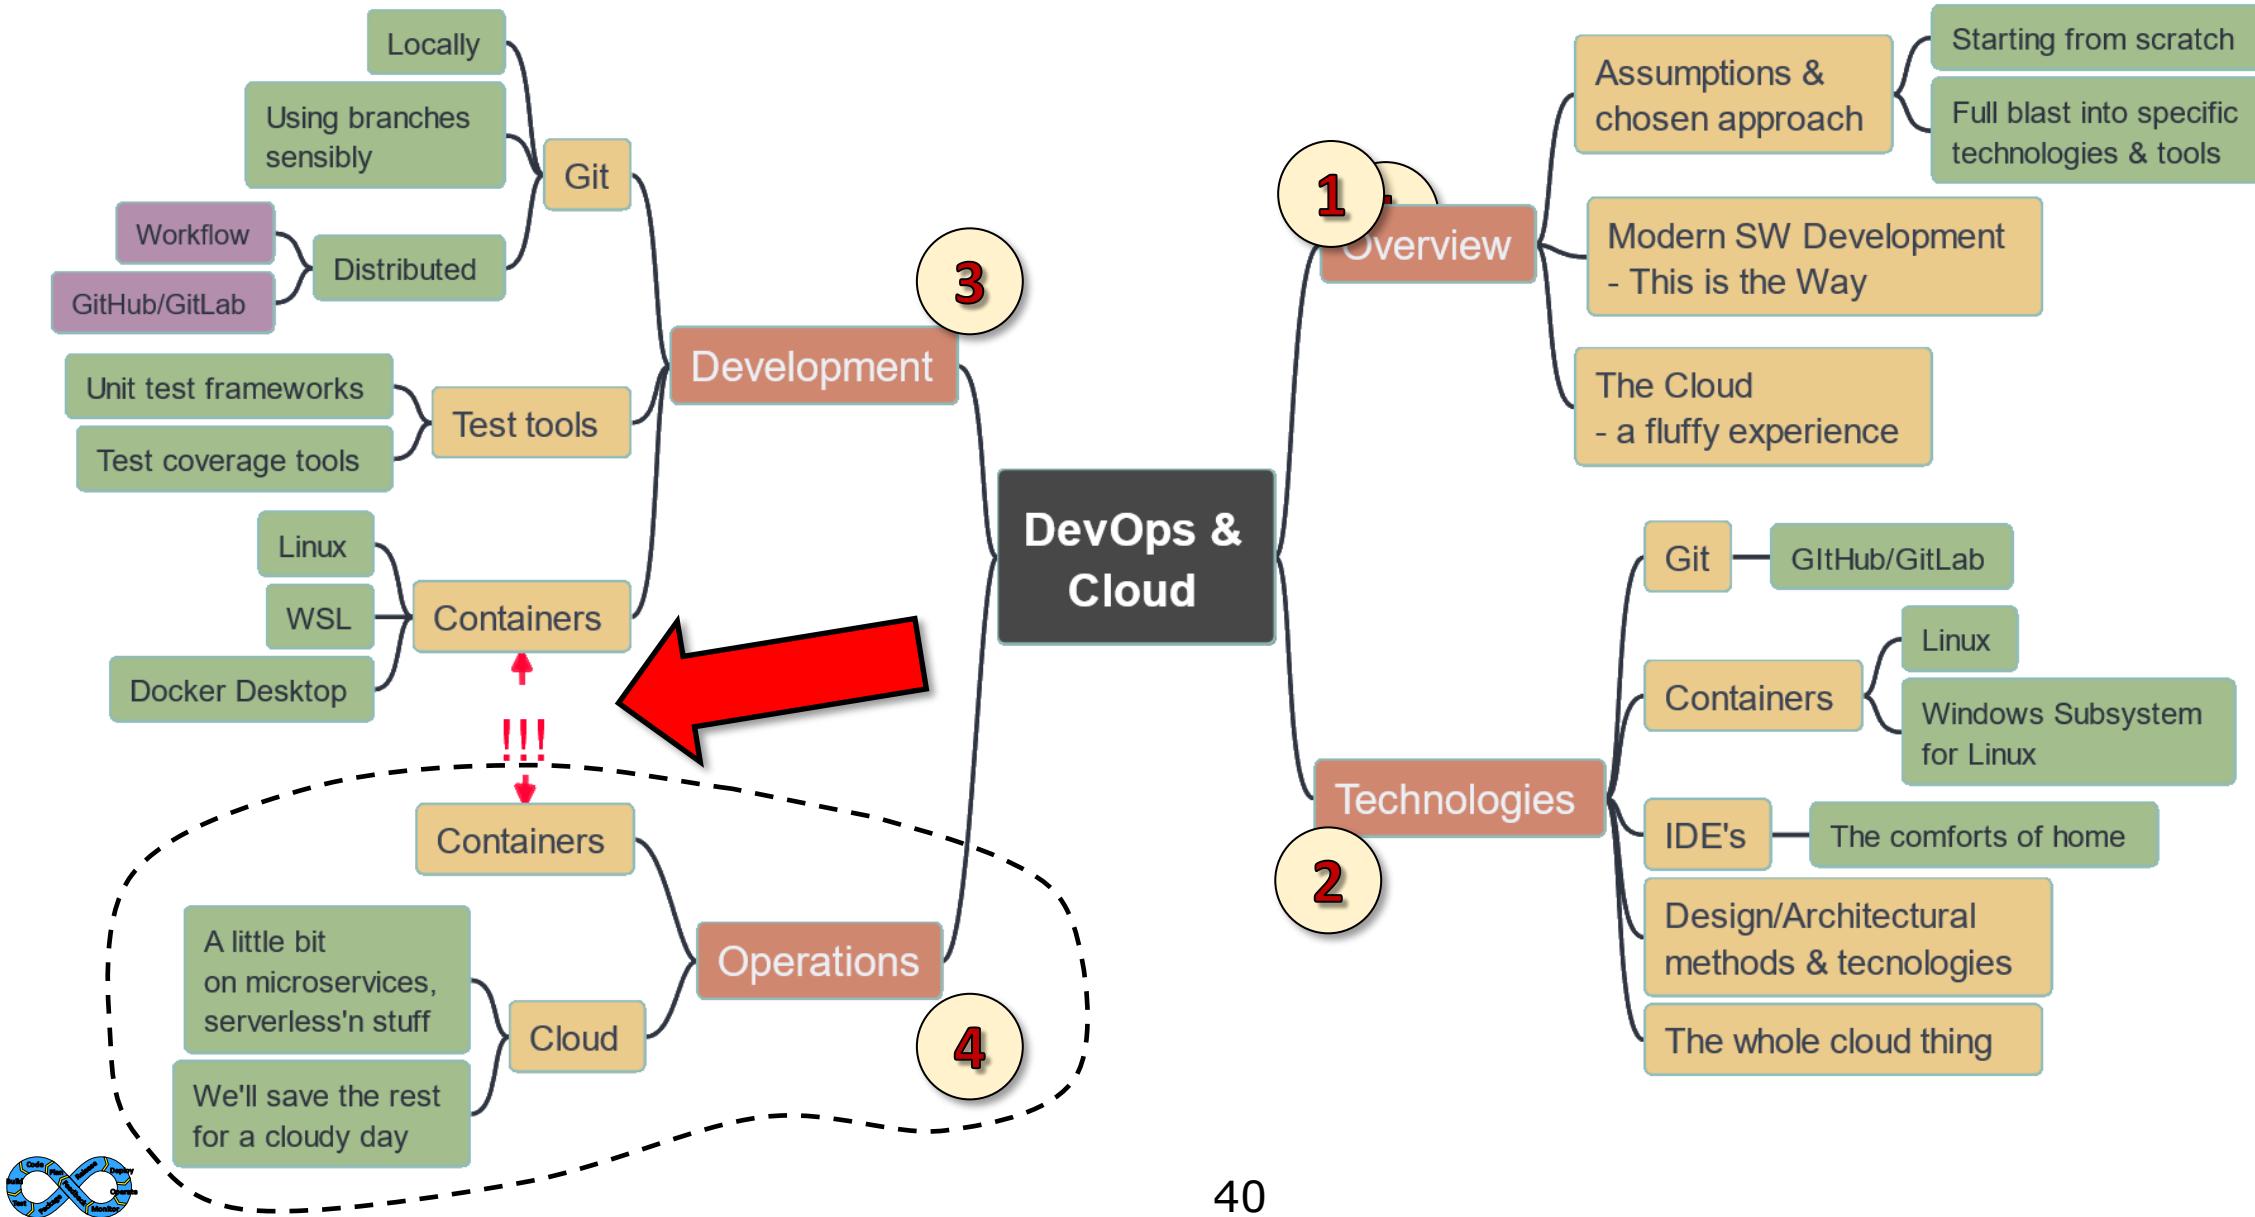


- WSL – bashing around .....
- Docker Desktop
- Working with a project using a *development* container  
<https://github.com/DaFessor/devops-demo>
  - Demo full development environment w. Java & Maven
  - Run application: mvn spring-boot:run
  - Run tests: mvn test
  - Show result of doing test coverage: <https://dafessor.github.io/devops-demo/>
- VSCode really rocks when it comes to container support





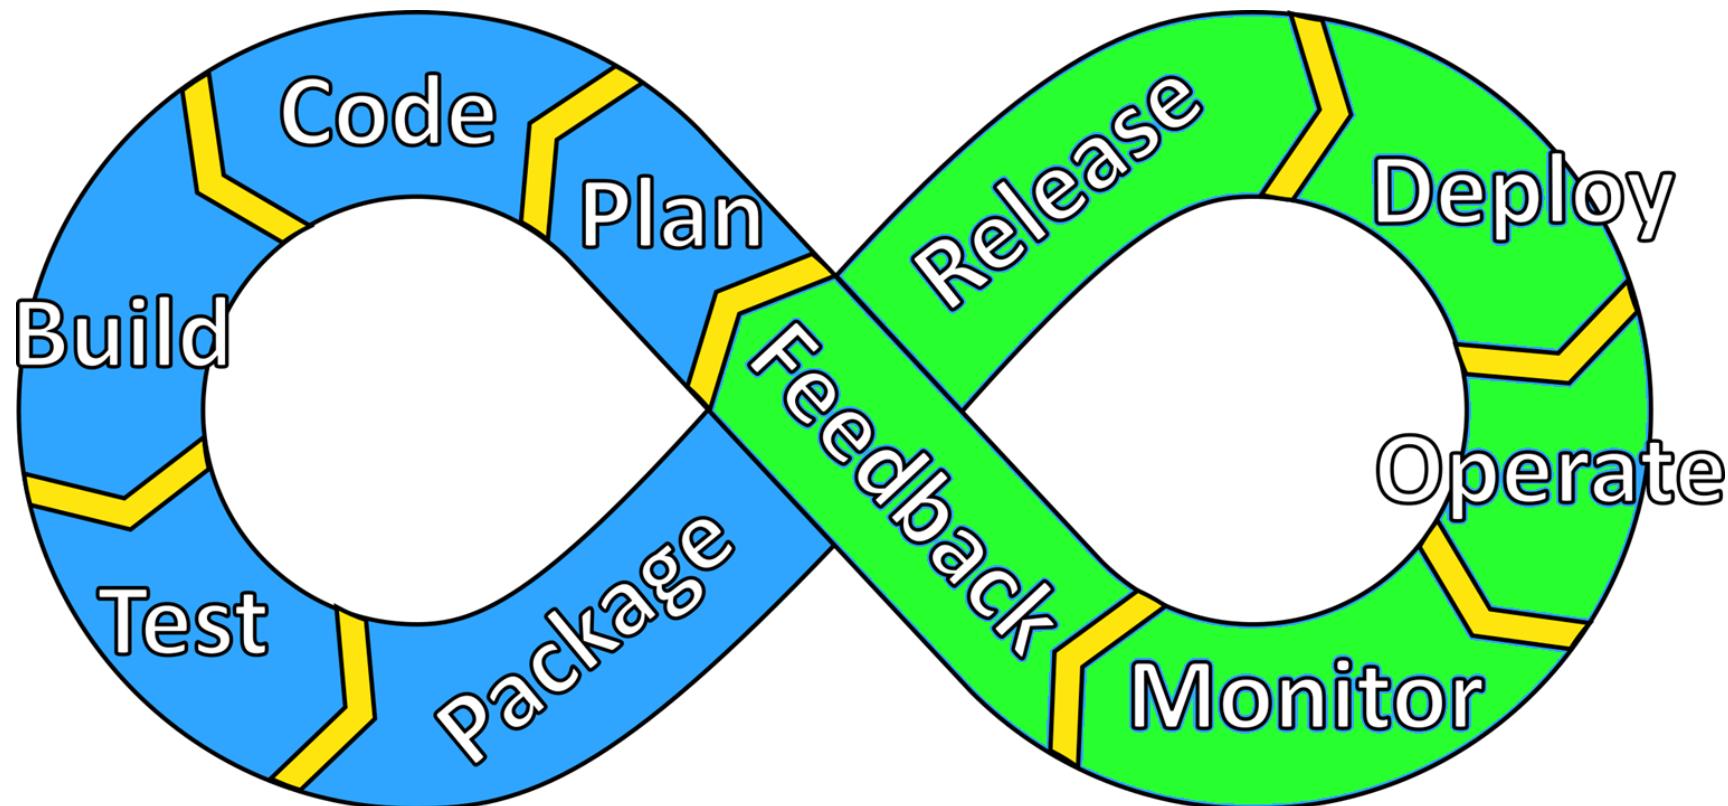
# Overview – mind this map





# The Ops-part of DevOps – tools & technologies

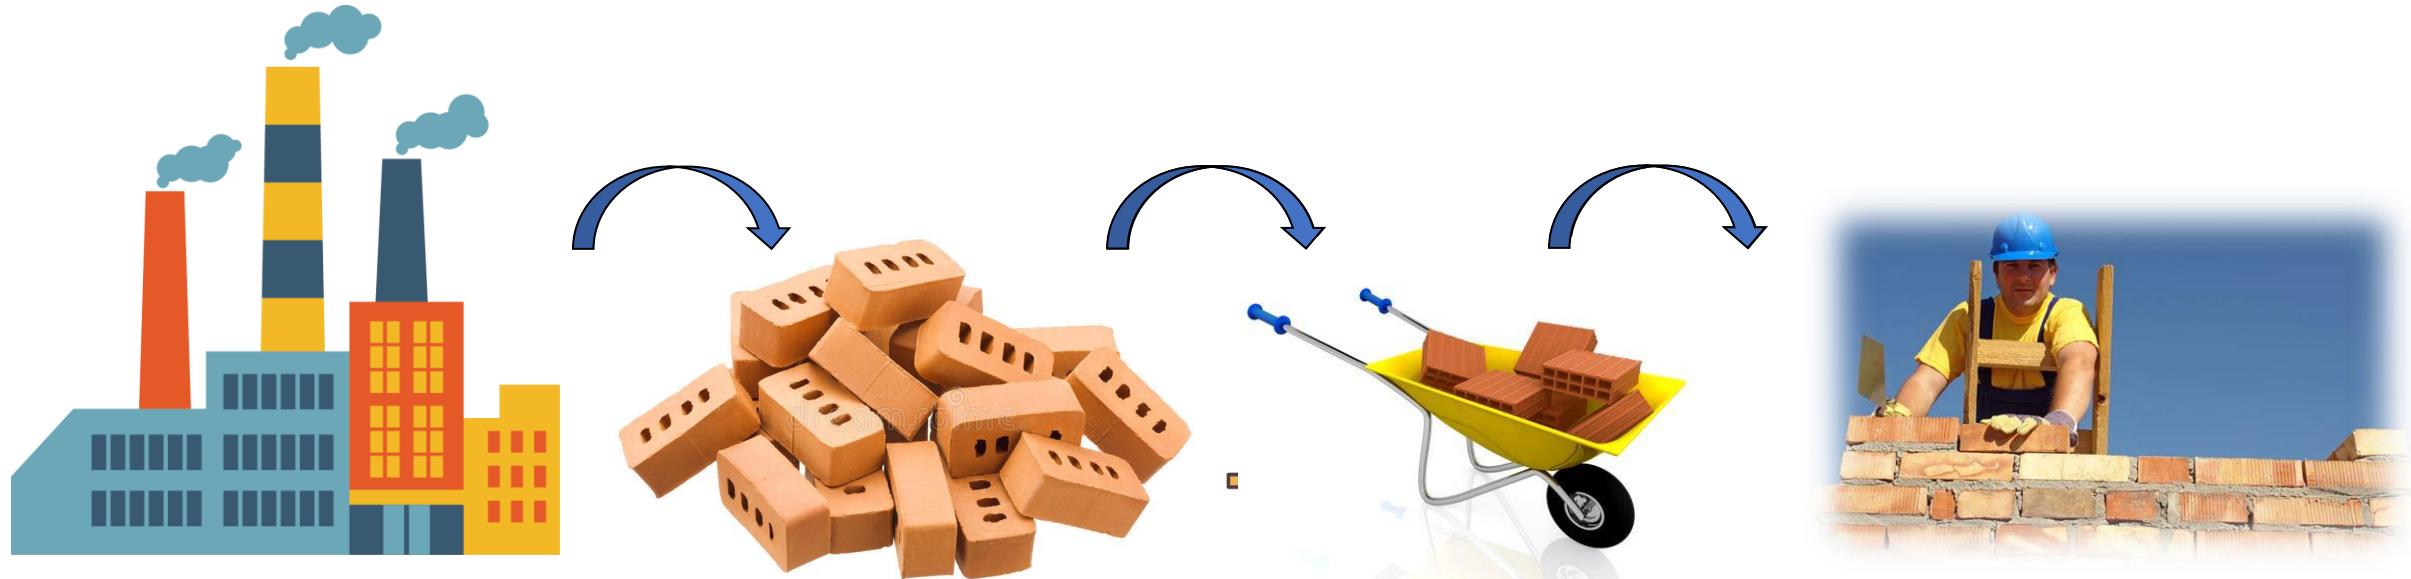
- Now let's take a look at (some of) the tools/technologies we can use for DevOps'ing the Ops part ....





# So, we get Dev-ilish fast at releasing new versions ...

- ... but that's not much help if we can't put the new software to use



- One of the corner stones of the DevOps “mind set” is that we have to think about the full sw lifecycle – so we’re not done just because we throw a new release out the door ....





# Facilitating the Ops-part of the sw lifecycle

- How can we make it easy to deploy and operate our software?
- Let's start with the easy/obvious thing - we need to deliver our software in an easy-to-deploy “format”, but what is that?
  - executable/jar-file(s)?
  - container(s)? <- This has become increasingly popular
- We also need to take operational demands into account
  - Is a stop-update-restart cycle permissible or do we have a 100% uptime requirement (think Netflix)?
- If we are not allowed any downtime for updates, we must design the software from the ground up to facilitate a “running” update – we need to think about how we design the basic architecture





# Updating software on the run

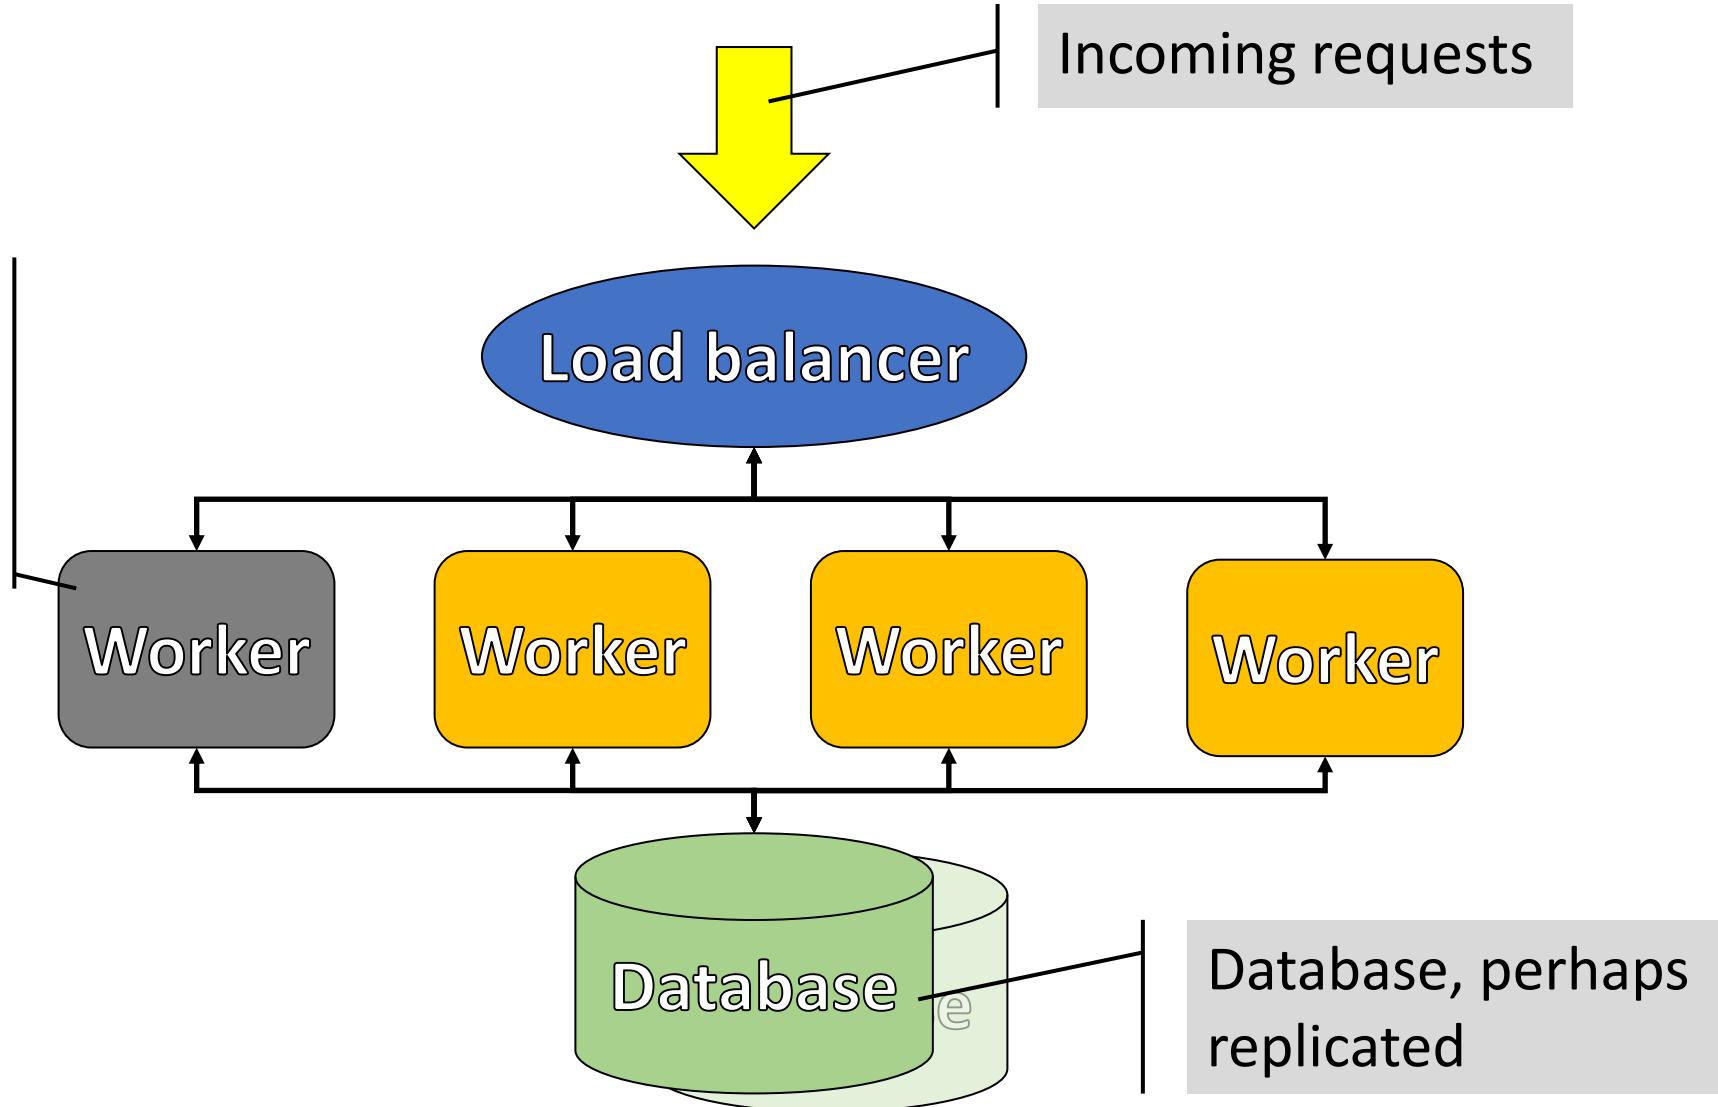
- This is a rather large subject, but let's take a quick look on some of the ways we may make it possible to do “hot updates”
- We'll look at
  - Load balancing
  - Micro services
  - Serverless





# Load balancing

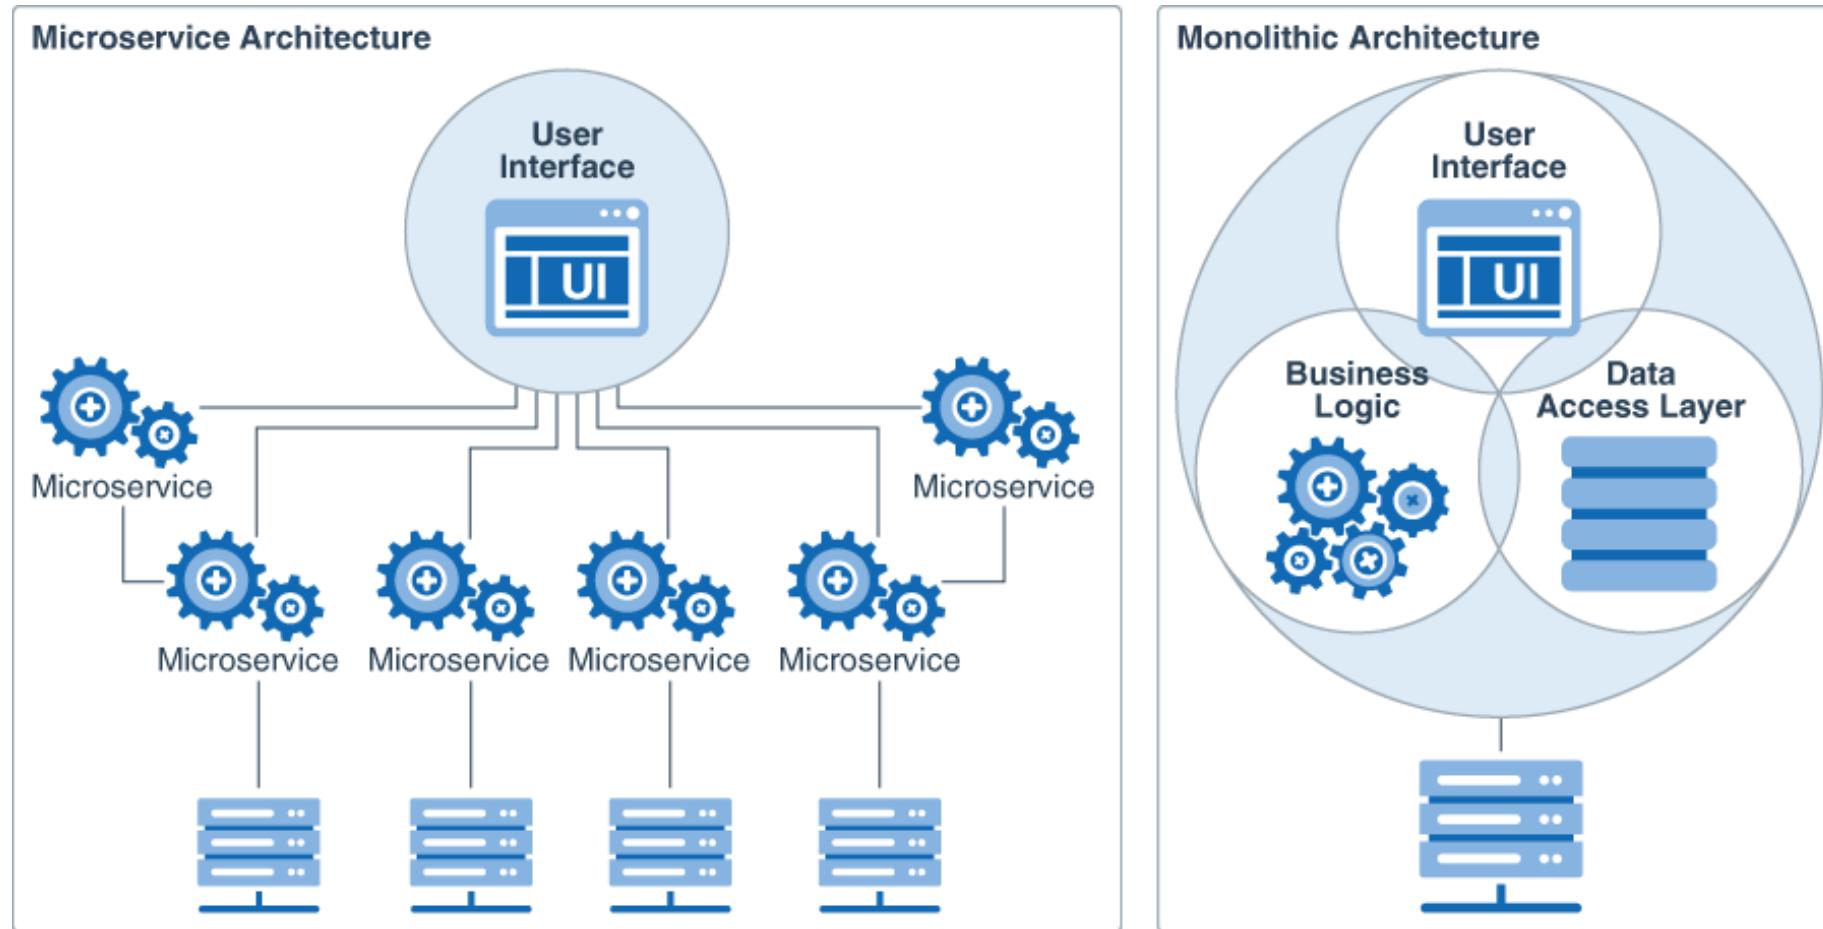
We can do a stop-update-restart cycle on 1 worker at a time w.o. interrupting normal operations





# Micro services (1)

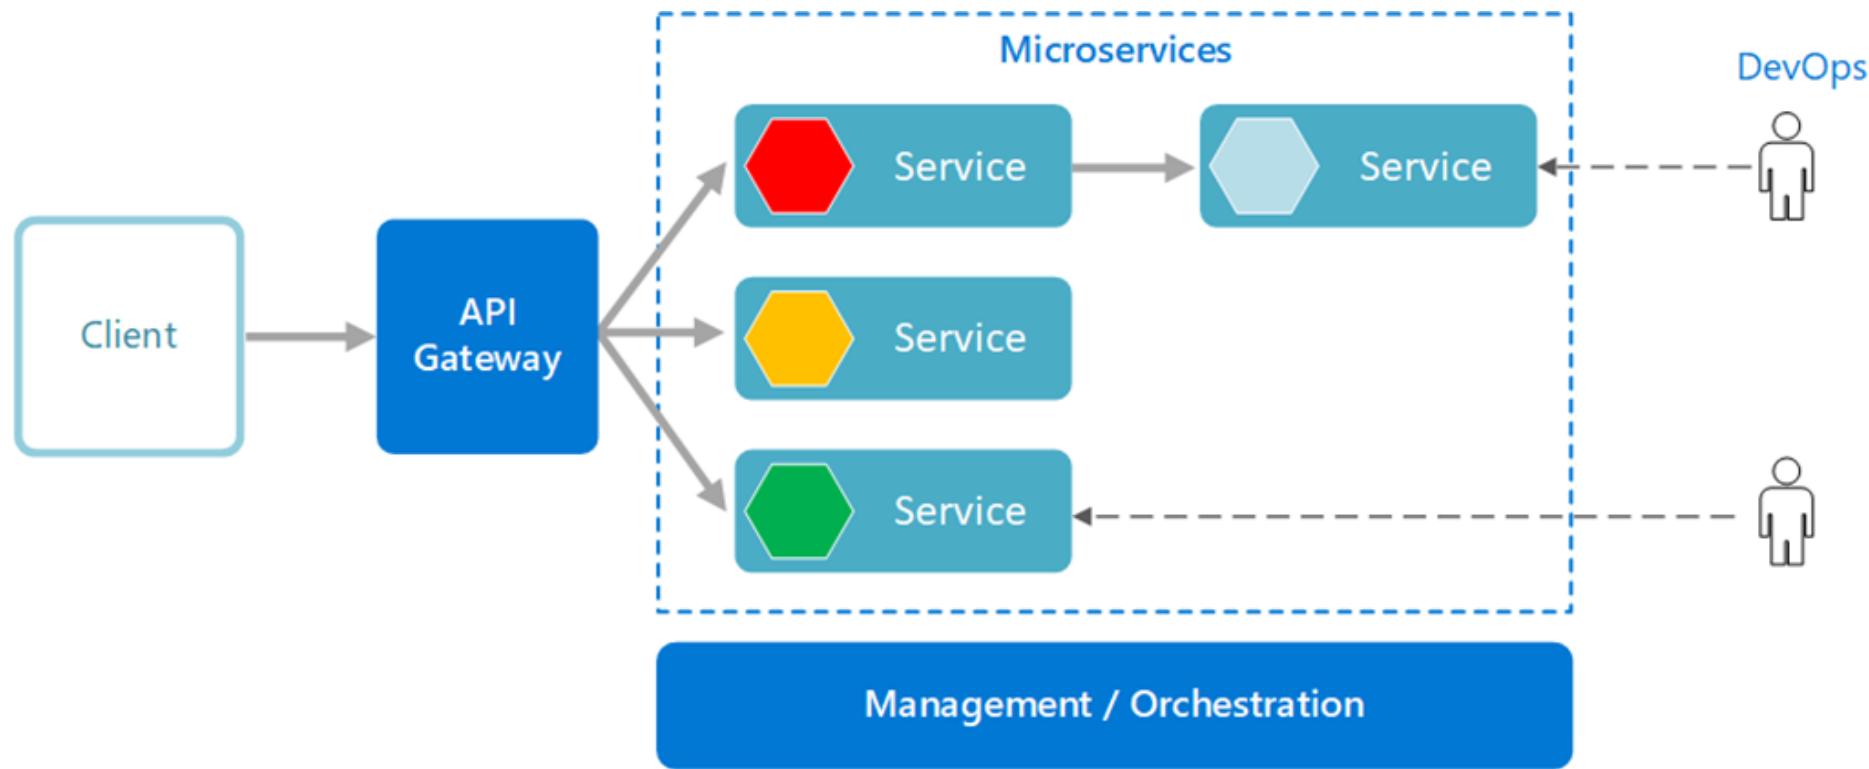
- Monolithic vs microservice architecture





# Microservices (2)

- A microservice architecture opens for a very flexible DevOps setup – you can have different team/tool/DevOps setup per microservice





# Microservices (3)

- Microservices frequently runs in containers in the Cloud
- But while microservices frequently is associated with running software in the Cloud, it's actually a general architectural pattern that can be used without necessarily using the Cloud (or perhaps with a private on-premise cloud)
- So, in many ways microservices is an architectural pattern that can be used to get rid of traditional monolithic applications with all their “bad qualities”
- A microservice can be developed, maintained and concurrently updated separately from all other parts of the system – that's smart!





# Serverless

- No, Serverless is not w.o servers – there has to be some servers running, somewhere
- Serverless is cloud-native – i.e. closely coupled to Cloud service providers
- Serverless is also known as FaaS = Function-as-a-Service
- Cloud providers sells ready-to-use serverless databases, serverless storage, serverless blahblah etc.
- But what IS serverless – it's actually hard to find a single, simple and direct answer on that ....





# Serverless explained (sort of)

- Imagine architecting your code as a set of functions
- Each time the code needs to do some kind of work, a specific function is called to do the job
- The code for the function is normally dormant, i.e. the code is not “active” like it would be if it was placed inside a traditional server application – so no CPU cycles are used to keep the code “warm” or ready
- Whenever the function is called, the code is spun up very fast, it gets its data, does what it needs, returns data and stops again
- You need to pay for the number of times you call the code, and the amount of CPU cycles that’s consumed during the call – it does not cost anything to have the function “on standby”, you only pay for actual use





# Interesting aspects of μ-services and serverless

- Some say that microservices and serverless has the potential to make the traditional Ops-work obsolete
- The idea is that developers themselves become Ops-people, if they get the responsibility of updating microservices or serverless functions directly whenever the developers make new versions available
- So, we don't anybody to babysit the operations and supervise updates of the software
- That has caused some people to declare that “normal” DevOps thinking will become obsolete, since the Dev part and Ops part fuse together
- Time will tell ...

